

ML_Finance_HW2

February 23, 2023

0.0.1 • Goal: Learn and use different models to solve a multi-class classification problem.

0.0.2 • Background: This data collects personal information of clients. The management wants to build an intelligent system to segregate the people into credit score brackets to reduce the manual efforts

0.0.3 • Author: Dhyey Dharmendrakumar Mavani

```
[1]: # checking the current working directory access
import os
print(os.getcwd())
```

```
/Users/dhyeymavani/Library/CloudStorage/GoogleDrive-dmavani25@amherst.edu/My
Drive/Columbia VUS/SPRING2023/MATH GR 5430 MACHINE LEARNING FOR
FINANCE/ML_Finance_HW2
```

1 1. Setup and Data Fetching (5 points)

```
[2]: # importing the data from the given csv (with cols separated by ";")
import pandas as pd
import numpy as np
client_data = pd.read_csv("./HW2.csv")
client_data
```

```
/var/folders/rb/x4vhwjb16wj764ztdb96gnlh0000gn/T/ipykernel_27457/3103777598.py:4
: DtypeWarning: Columns (26) have mixed types. Specify dtype option on import or
set low_memory=False.
```

```
client_data = pd.read_csv("./HW2.csv")
```

```
[2]:
```

	ID	Customer_ID	Month	Name	Age	SSN	\
0	0x1602	CUS_0xd40	January	Aaron Maashoh	23	821-00-0265	
1	0x1603	CUS_0xd40	February	Aaron Maashoh	23	821-00-0265	
2	0x1604	CUS_0xd40	March	Aaron Maashoh	-500	821-00-0265	
3	0x1605	CUS_0xd40	April	Aaron Maashoh	23	821-00-0265	
4	0x1606	CUS_0xd40	May	Aaron Maashoh	23	821-00-0265	
...	
99995	0x25fe9	CUS_0x942c	April	Nicks	25	078-73-5990	

99996	0x25fea	CUS_0x942c	May	Nicks	25	078-73-5990
99997	0x25feb	CUS_0x942c	June	Nicks	25	078-73-5990
99998	0x25fec	CUS_0x942c	July	Nicks	25	078-73-5990
99999	0x25fed	CUS_0x942c	August	Nicks	25	078-73-5990

	Occupation	Annual_Income	Monthly_Inhand_Salary	Num_Bank_Accounts	...	\
0	Scientist	19114.12	1824.843333	3	...	
1	Scientist	19114.12	NaN	3	...	
2	Scientist	19114.12	NaN	3	...	
3	Scientist	19114.12	NaN	3	...	
4	Scientist	19114.12	1824.843333	3	...	
...	
99995	Mechanic	39628.99	3359.415833	4	...	
99996	Mechanic	39628.99	3359.415833	4	...	
99997	Mechanic	39628.99	3359.415833	4	...	
99998	Mechanic	39628.99	3359.415833	4	...	
99999	Mechanic	39628.99_	3359.415833	4	...	

	Credit_Mix	Outstanding_Debt	Credit_Utilization_Ratio	\
0	-	809.98	26.822620	
1	Good	809.98	31.944960	
2	Good	809.98	28.609352	
3	Good	809.98	31.377862	
4	Good	809.98	24.797347	
...	
99995	-	502.38	34.663572	
99996	-	502.38	40.565631	
99997	Good	502.38	41.255522	
99998	Good	502.38	33.638208	
99999	Good	502.38	34.192463	

	Credit_History_Age	Payment_of_Min_Amount	Total_EMI_per_month	\
0	22 Years and 1 Months	No	49.574949	
1	NaN	No	49.574949	
2	22 Years and 3 Months	No	49.574949	
3	22 Years and 4 Months	No	49.574949	
4	22 Years and 5 Months	No	49.574949	
...	
99995	31 Years and 6 Months	No	35.104023	
99996	31 Years and 7 Months	No	35.104023	
99997	31 Years and 8 Months	No	35.104023	
99998	31 Years and 9 Months	No	35.104023	
99999	31 Years and 10 Months	No	35.104023	

	Amount_invested_monthly	Payment_Behaviour	\
0	80.41529543900253	High_spent_Small_value_payments	
1	118.28022162236736	Low_spent_Large_value_payments	

2	81.699521264648	Low_spent_Medium_value_payments
3	199.4580743910713	Low_spent_Small_value_payments
4	41.420153086217326	High_spent_Medium_value_payments
...
99995	60.97133255718485	High_spent_Large_value_payments
99996	54.18595028760385	High_spent_Medium_value_payments
99997	24.02847744864441	High_spent_Large_value_payments
99998	251.67258219721603	Low_spent_Large_value_payments
99999	167.1638651610451	!@9#%8

	Monthly_Balance	Credit_Score
0	312.49408867943663	Good
1	284.62916249607184	Good
2	331.2098628537912	Good
3	223.45130972736786	Good
4	341.48923103222177	Good
...
99995	479.866228	Poor
99996	496.65161	Poor
99997	516.809083	Poor
99998	319.164979	Standard
99999	393.673696	Poor

[100000 rows x 28 columns]

2. Exploratory Data Analysis (20 points)

- **Data cleaning:** Some features have missing or invalid values. Choose the way you see as appropriate to clean the dataframe so that there are no missing or invalid values.

```
[3]: client_data_after_dropping_nas = client_data.dropna()
      client_data_after_dropping_nas.reset_index(inplace = True)

      for i in range(len(client_data_after_dropping_nas['Age'])):
          alpha_num_str = client_data_after_dropping_nas['Age'][i]
          num_str = ''.join((z for z in alpha_num_str if z.isdigit()))
          client_data_after_dropping_nas.iat[i, 5] = num_str

      client_data_after_dropping_nas["Age"] = client_data_after_dropping_nas["Age"].
          ↪astype(int)
      temp_age_filtered_client_data =
          ↪client_data_after_dropping_nas[client_data_after_dropping_nas['Age'] < 100]
      age_filtered_client_data =
          ↪temp_age_filtered_client_data[temp_age_filtered_client_data['Age'] > 0]

      occupation_filtered_client_data =
          ↪age_filtered_client_data[age_filtered_client_data['Occupation'] != "_____"]
```

```

del occupation_filtered_client_data['index']

occupation_filtered_client_data = occupation_filtered_client_data.reset_index()

for i in range(len(occupation_filtered_client_data['Annual_Income'])):
    alpha_num_str = occupation_filtered_client_data['Annual_Income'][i]
    num_str = ''.join((z for z in alpha_num_str if (z.isdigit() or z == ".")))
    head, sep, tail = num_str.partition('.')
    occupation_filtered_client_data.iat[i, 8] = head

occupation_filtered_client_data['Annual_Income'] =
    ↳occupation_filtered_client_data['Annual_Income'].astype(int)

del occupation_filtered_client_data['index']

occupation_filtered_client_data = occupation_filtered_client_data.reset_index()

for i in range(len(occupation_filtered_client_data['Num_of_Loan'])):
    alpha_num_str = occupation_filtered_client_data['Num_of_Loan'][i]
    num_str = ''.join((z for z in alpha_num_str if (z.isdigit())))
    head, sep, tail = num_str.partition('_')
    occupation_filtered_client_data.iat[i, 13] = num_str

occupation_filtered_client_data["Num_of_Loan"] =
    ↳occupation_filtered_client_data["Num_of_Loan"].astype(int)

loan_filtered_client_data =
    ↳occupation_filtered_client_data[occupation_filtered_client_data["Num_of_Loan"]
    ↳>= 0]
loan_filtered_client_data =
    ↳loan_filtered_client_data[loan_filtered_client_data["Num_of_Loan"] <= 10]

bank_acc_filtered_client_data =
    ↳loan_filtered_client_data[loan_filtered_client_data["Num_Bank_Accounts"] >=
    ↳0]
bank_acc_filtered_client_data =
    ↳bank_acc_filtered_client_data[bank_acc_filtered_client_data["Num_Bank_Accounts"]
    ↳<= 10]

credit_card_filtered_client_data =
    ↳bank_acc_filtered_client_data[bank_acc_filtered_client_data["Num_Credit_Card"]
    ↳>= 0]
credit_card_filtered_client_data =
    ↳credit_card_filtered_client_data[credit_card_filtered_client_data["Num_Credit_Card"]
    ↳<= 10]

```

```

interest_filtered_client_data =
    ↳ credit_card_filtered_client_data[credit_card_filtered_client_data["Interest_Rate"]
    ↳ >= 0]
interest_filtered_client_data =
    ↳ interest_filtered_client_data[interest_filtered_client_data["Interest_Rate"]
    ↳ <= 35]

delay_filtered_client_data =
    ↳ interest_filtered_client_data[interest_filtered_client_data["Delay_from_due_date"]
    ↳ >= 0]
delay_filtered_client_data =
    ↳ delay_filtered_client_data[delay_filtered_client_data["Delay_from_due_date"]
    ↳ <= 65]

del delay_filtered_client_data['index']

delay_filtered_client_data = delay_filtered_client_data.reset_index()

for i in range(len(delay_filtered_client_data["Num_of_Delayed_Payment"])):
    alpha_num_str = delay_filtered_client_data["Num_of_Delayed_Payment"][i]
    num_str = ''.join((z for z in alpha_num_str if (z.isdigit())))
    head, sep, tail = num_str.partition('_')
    delay_filtered_client_data.iat[i, 16] = head

delay_filtered_client_data["Num_of_Delayed_Payment"] =
    ↳ delay_filtered_client_data["Num_of_Delayed_Payment"].astype(int)

delayed_pay_filtered_client_data =
    ↳ delay_filtered_client_data[delay_filtered_client_data["Num_of_Delayed_Payment"]
    ↳ >= 0]
delayed_pay_filtered_client_data =
    ↳ delayed_pay_filtered_client_data[delayed_pay_filtered_client_data["Num_of_Delayed_Payment"]
    ↳ <= 25]

del delayed_pay_filtered_client_data['index']

delayed_pay_filtered_client_data = delayed_pay_filtered_client_data.
    ↳ reset_index()
delayed_pay_filtered_client_data.replace("_", np.nan, inplace=True)
delayed_pay_filtered_client_data.dropna(inplace=True)

delayed_pay_filtered_client_data["Changed_Credit_Limit"] =
    ↳ delayed_pay_filtered_client_data["Changed_Credit_Limit"].astype(float)

```

```

cred_inq_filtered_client_data =
    ↪ delayed_pay_filtered_client_data[delayed_pay_filtered_client_data["Num_Credit_Inquiries"]
    ↪ >= 0]

cred_inq_filtered_client_data =
    ↪ cred_inq_filtered_client_data[cred_inq_filtered_client_data["Num_Credit_Inquiries"]
    ↪ <= 25]

credit_mix_filtered_client_data =
    ↪ cred_inq_filtered_client_data[cred_inq_filtered_client_data['Credit_Mix'] !=
    ↪ "_"]

del credit_mix_filtered_client_data['index']

credit_mix_filtered_client_data = credit_mix_filtered_client_data.reset_index()

for i in range(len(credit_mix_filtered_client_data['Outstanding_Debt'])):
    alpha_num_str = credit_mix_filtered_client_data['Outstanding_Debt'][i]
    num_str = ''.join((z for z in alpha_num_str if (z.isdigit() or z == ".")))
    head, sep, tail = num_str.partition('_')
    credit_mix_filtered_client_data.iat[i, 20] = head

credit_mix_filtered_client_data['Outstanding_Debt'] =
    ↪ credit_mix_filtered_client_data['Outstanding_Debt'].astype(float)

del credit_mix_filtered_client_data['index']

credit_mix_filtered_client_data = credit_mix_filtered_client_data.reset_index()

for i in range(len(credit_mix_filtered_client_data['Credit_History_Age'])):
    alpha_num_str = credit_mix_filtered_client_data['Credit_History_Age'][i]
    res = [float(i) for i in alpha_num_str.split() if i.isdigit()]
    y, m = res
    age = (y + m/12)
    credit_mix_filtered_client_data.iat[i, 22] = age

credit_mix_filtered_client_data['Credit_History_Age'] =
    ↪ credit_mix_filtered_client_data['Credit_History_Age'].astype(float)

credit_mix_filtered_client_data.replace("NM", np.nan, inplace=True)
credit_mix_filtered_client_data.dropna(inplace=True)

credit_mix_filtered_client_data.replace("__10000__", "10000.00", inplace=True)
credit_mix_filtered_client_data['Amount_invested_monthly'] =
    ↪ credit_mix_filtered_client_data['Amount_invested_monthly'].astype(float)

credit_mix_filtered_client_data.replace("@9#%8", np.nan, inplace=True)

```

```

credit_mix_filtered_client_data.dropna(inplace=True)

credit_mix_filtered_client_data.replace("__-333333333333333333333333__", np.
    ↳nan, inplace=True)
credit_mix_filtered_client_data.dropna(inplace=True)
credit_mix_filtered_client_data['Monthly_Balance'] =_
    ↳credit_mix_filtered_client_data['Monthly_Balance'].astype(float)

cleaned_client_data = credit_mix_filtered_client_data.copy()

del cleaned_client_data['index']

cleaned_client_data = cleaned_client_data.reset_index()

del cleaned_client_data['index']

cleaned_client_data

```

/var/folders/rb/x4vhwjb16wj764ztdb96gnlh0000gn/T/ipykernel_27457/3187363897.py:9
: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
client_data_after_dropping_nas["Age"] =
client_data_after_dropping_nas["Age"].astype(int)

```

[3]:
      ID Customer_ID      Month      Name  Age      SSN  \
0    0x1608  CUS_0xd40      July  Aaron Maashoh  23  821-00-0265
1    0x160f  CUS_0x21b1  February  Rick Rothackerj  28  004-07-5839
2    0x1612  CUS_0x21b1      May  Rick Rothackerj  28  004-07-5839
3    0x1613  CUS_0x21b1      June  Rick Rothackerj  28  004-07-5839
4    0x1615  CUS_0x21b1  August  Rick Rothackerj  28  004-07-5839
...    ...    ...    ...    ...    ...    ...
26523  0x25fb6  CUS_0x372c  January  Lucia Mutikanik  18  340-85-7301
26524  0x25fce  CUS_0xaf61  January  Chris Wickhamm  49  133-16-7738
26525  0x25fcf  CUS_0xaf61  February  Chris Wickhamm  49  133-16-7738
26526  0x25fdb  CUS_0x8600  February  Sarah McBridec  28  031-35-0942
26527  0x25fe1  CUS_0x8600  August  Sarah McBridec  29  031-35-0942

      Occupation  Annual_Income  Monthly_Inhand_Salary  Num_Bank_Accounts  \
0    Scientist      19114      1824.843333      3
1    Teacher      34847      3037.986667      2
2    Teacher      34847      3037.986667      2
3    Teacher      34847      3037.986667      2
4    Teacher      34847      3037.986667      2

```

...
26523	Lawyer	42903	3468.315833	0
26524	Writer	37188	3097.008333	1
26525	Writer	37188	3097.008333	1
26526	Architect	20002	1929.906667	10
26527	Architect	20002	1929.906667	10

	...	Credit_Mix	Outstanding_Debt	Credit_Utilization_Ratio	\
0	...	Good	809.98	22.537593	
1	...	Good	605.03	38.550848	
2	...	Good	605.03	34.977895	
3	...	Good	605.03	33.381010	
4	...	Good	605.03	32.933856	

...
26523	...	Good	1079.48	27.289440
26524	...	Good	620.64	39.080823
26525	...	Good	620.64	32.803431
26526	...	Bad	3571.70	39.772607
26527	...	Bad	3571.70	37.140784

	Credit_History_Age	Payment_of_Min_Amount	Total_EMI_per_month	\
0	22.583333	No	49.574949	
1	26.666667	No	18.816215	
2	26.916667	No	18.816215	
3	27.000000	No	18.816215	
4	27.166667	No	18.816215	

...
26523	28.083333	No	50894.000000
26524	29.750000	No	84.205949
26525	29.833333	No	84.205949
26526	5.750000	Yes	12112.000000
26527	6.250000	Yes	60.964772

	Amount_invested_monthly	Payment_Behaviour	\
0	178.344067	Low_spent_Small_value_payments	
1	40.391238	High_spent_Large_value_payments	
2	130.115420	Low_spent_Small_value_payments	
3	43.477190	High_spent_Large_value_payments	
4	218.904344	Low_spent_Small_value_payments	

...
26523	78.514945	High_spent_Small_value_payments
26524	223.875018	Low_spent_Small_value_payments
26525	70.869970	High_spent_Large_value_payments
26526	148.275233	Low_spent_Small_value_payments
26527	34.662906	High_spent_Large_value_payments

Monthly_Balance Credit_Score

0	244.565317	Good
1	484.591214	Good
2	444.867032	Good
3	481.505262	Good
4	356.078109	Good
...
26523	493.341182	Good
26524	291.619866	Good
26525	394.624914	Good
26526	273.750662	Poor
26527	337.362988	Standard

[26528 rows x 28 columns]

```
[4]: cleaned_client_data.to_csv("clean_client_data.csv")
```

Columns in the cleaned data:

- 'Unnamed: 0'
- 'ID'
- 'Customer_ID'
- 'Month'
- 'Name'
- 'Age'
- 'SSN'
- 'Occupation'
- 'Annual_Income'
- 'Monthly_Inhand_Salary'
- 'Num_Bank_Accounts'
- 'Num_Credit_Card'
- 'Interest_Rate'
- 'Num_of_Loan'
- 'Type_of_Loan'
- 'Delay_from_due_date'
- 'Num_of_Delayed_Payment'
- 'Changed_Credit_Limit'
- 'Num_Credit_Inquiries'
- 'Credit_Mix'
- 'Outstanding_Debt'
- 'Credit_Utilization_Ratio'
- 'Credit_History_Age'
- 'Payment_of_Min_Amount'
- 'Total_EMI_per_month'
- 'Amount_invested_monthly'
- 'Payment_Behaviour'
- 'Monthly_Balance'
- 'Credit_Score'

```
[5]: data = pd.read_csv("./clean_client_data.csv")
del data["Unnamed: 0"]
data
```

```
[5]:
```

	ID	Customer_ID	Month	Name	Age	SSN	\
0	0x1608	CUS_0xd40	July	Aaron Maashoh	23	821-00-0265	
1	0x160f	CUS_0x21b1	February	Rick Rothackerj	28	004-07-5839	
2	0x1612	CUS_0x21b1	May	Rick Rothackerj	28	004-07-5839	
3	0x1613	CUS_0x21b1	June	Rick Rothackerj	28	004-07-5839	
4	0x1615	CUS_0x21b1	August	Rick Rothackerj	28	004-07-5839	
...	
26523	0x25fb6	CUS_0x372c	January	Lucia Mutikanik	18	340-85-7301	
26524	0x25fce	CUS_0xaf61	January	Chris Wickhamm	49	133-16-7738	
26525	0x25fcf	CUS_0xaf61	February	Chris Wickhamm	49	133-16-7738	
26526	0x25fdb	CUS_0x8600	February	Sarah McBridec	28	031-35-0942	
26527	0x25fe1	CUS_0x8600	August	Sarah McBridec	29	031-35-0942	

	Occupation	Annual_Income	Monthly_Inhand_Salary	Num_Bank_Accounts	\
0	Scientist	19114	1824.843333	3	
1	Teacher	34847	3037.986667	2	
2	Teacher	34847	3037.986667	2	
3	Teacher	34847	3037.986667	2	
4	Teacher	34847	3037.986667	2	
...	
26523	Lawyer	42903	3468.315833	0	
26524	Writer	37188	3097.008333	1	
26525	Writer	37188	3097.008333	1	
26526	Architect	20002	1929.906667	10	
26527	Architect	20002	1929.906667	10	

	...	Credit_Mix	Outstanding_Debt	Credit_Utilization_Ratio	\
0	...	Good	809.98	22.537593	
1	...	Good	605.03	38.550848	
2	...	Good	605.03	34.977895	
3	...	Good	605.03	33.381010	
4	...	Good	605.03	32.933856	
...	
26523	...	Good	1079.48	27.289440	
26524	...	Good	620.64	39.080823	
26525	...	Good	620.64	32.803431	
26526	...	Bad	3571.70	39.772607	
26527	...	Bad	3571.70	37.140784	

	Credit_History_Age	Payment_of_Min_Amount	Total_EMI_per_month	\
0	22.583333	No	49.574949	
1	26.666667	No	18.816215	
2	26.916667	No	18.816215	

3	27.000000	No	18.816215
4	27.166667	No	18.816215
...
26523	28.083333	No	50894.000000
26524	29.750000	No	84.205949
26525	29.833333	No	84.205949
26526	5.750000	Yes	12112.000000
26527	6.250000	Yes	60.964772

	Amount_invested_monthly	Payment_Behaviour \
0	178.344067	Low_spent_Small_value_payments
1	40.391238	High_spent_Large_value_payments
2	130.115420	Low_spent_Small_value_payments
3	43.477190	High_spent_Large_value_payments
4	218.904344	Low_spent_Small_value_payments
...
26523	78.514945	High_spent_Small_value_payments
26524	223.875018	Low_spent_Small_value_payments
26525	70.869970	High_spent_Large_value_payments
26526	148.275233	Low_spent_Small_value_payments
26527	34.662906	High_spent_Large_value_payments


	Monthly_Balance	Credit_Score
0	244.565317	Good
1	484.591214	Good
2	444.867032	Good
3	481.505262	Good
4	356.078109	Good
...
26523	493.341182	Good
26524	291.619866	Good
26525	394.624914	Good
26526	273.750662	Poor
26527	337.362988	Standard


[26528 rows x 28 columns]

- Plot the distributions of two numerical features of your choice. What do you discover about those variables?

```
[6]: import seaborn as sns
sns.displot(data, x="Age", kind="kde")
sns.displot(data, x="Annual_Income", kind="kde")
```

```
[6]: <seaborn.axisgrid.FacetGrid at 0x7fef9c445d60>
```






We have chosen here to plot the density distributions of numeric variables age and annual income in our filtered dataset. We can see that both the distributions are right-skewed, while that of annual income is more so. Furthermore, we can see that the age with most density is around three peaks at 25, 38 and 42 approximately, while the density of annual income variable peaks when annual income takes a value of 0.05×10^7 (around 50,000). We can also see that age mostly lies between 10 to 60, while campaign values mostly lie between 0 to 1,000,000.

```
[36]: sns.scatterplot(data=data, x="Age", y="Annual_Income")
```

```
[36]: <AxesSubplot: xlabel='Age', ylabel='Annual_Income'>
```



Furthermore, on plotting the numeric variables of interest as a scatterplot with annual income on the y-axis and age on x-axis as shown above, we can see that annual income has almost the same seemingly random variation among different ages from 15 to 60 with most people in the low annual income range.

- **Plot the correlation matrix of all numeric features. What do you discover?**

```
[7]: data.corr()
```

```
[7]:
```

	Age	Annual_Income	Monthly_Inhand_Salary \
Age	1.000000	0.009402	0.085716
Annual_Income	0.009402	1.000000	0.035314
Monthly_Inhand_Salary	0.085716	0.035314	1.000000
Num_Bank_Accounts	-0.191913	-0.008823	-0.281451
Num_Credit_Card	-0.150085	-0.004909	-0.222464
Interest_Rate	-0.217417	-0.019732	-0.308845
Num_of_Loan	-0.196767	-0.002397	-0.249267
Delay_from_due_date	-0.170238	-0.011910	-0.246699
Num_of_Delayed_Payment	-0.182203	-0.004452	-0.282958
Changed_Credit_Limit	-0.154947	-0.008059	-0.173013
Num_Credit_Inquiries	-0.243182	-0.013630	-0.285780

Outstanding_Debt	-0.200344	-0.008690	-0.277441
Credit_Utilization_Ratio	0.022428	0.009059	0.159732
Credit_History_Age	0.232788	0.003094	0.283447
Total_EMI_per_month	0.000514	-0.010581	0.010987
Amount_invested_monthly	-0.004075	-0.000604	0.049266
Monthly_Balance	0.108362	0.020047	0.680436

	Num_Bank_Accounts	Num_Credit_Card	Interest_Rate \
Age	-0.191913	-0.150085	-0.217417
Annual_Income	-0.008823	-0.004909	-0.019732
Monthly_Inhand_Salary	-0.281451	-0.222464	-0.308845
Num_Bank_Accounts	1.000000	0.440339	0.579838
Num_Credit_Card	0.440339	1.000000	0.499532
Interest_Rate	0.579838	0.499532	1.000000
Num_of_Loan	0.469219	0.417180	0.539941
Delay_from_due_date	0.556531	0.481129	0.584251
Num_of_Delayed_Payment	0.595783	0.413135	0.561430
Changed_Credit_Limit	0.323093	0.245925	0.355906
Num_Credit_Inquiries	0.526000	0.468567	0.633609
Outstanding_Debt	0.515963	0.498871	0.633570
Credit_Utilization_Ratio	-0.060204	-0.048869	-0.066029
Credit_History_Age	-0.487774	-0.419828	-0.576757
Total_EMI_per_month	0.005370	-0.001450	0.003615
Amount_invested_monthly	-0.013920	-0.001580	-0.010806
Monthly_Balance	-0.284251	-0.235929	-0.318435

	Num_of_Loan	Delay_from_due_date \
Age	-0.196767	-0.170238
Annual_Income	-0.002397	-0.011910
Monthly_Inhand_Salary	-0.249267	-0.246699
Num_Bank_Accounts	0.469219	0.556531
Num_Credit_Card	0.417180	0.481129
Interest_Rate	0.539941	0.584251
Num_of_Loan	1.000000	0.490506
Delay_from_due_date	0.490506	1.000000
Num_of_Delayed_Payment	0.463587	0.537229
Changed_Credit_Limit	0.367170	0.286690
Num_Credit_Inquiries	0.544718	0.539331
Outstanding_Debt	0.640329	0.578490
Credit_Utilization_Ratio	-0.077027	-0.054367
Credit_History_Age	-0.596080	-0.487857
Total_EMI_per_month	-0.000914	-0.005720
Amount_invested_monthly	-0.019284	-0.004256
Monthly_Balance	-0.406558	-0.267302

	Num_of_Delayed_Payment	Changed_Credit_Limit \
Age	-0.182203	-0.154947

Annual_Income	-0.004452	-0.008059
Monthly_Inhand_Salary	-0.282958	-0.173013
Num_Bank_Accounts	0.595783	0.323093
Num_Credit_Card	0.413135	0.245925
Interest_Rate	0.561430	0.355906
Num_of_Loan	0.463587	0.367170
Delay_from_due_date	0.537229	0.286690
Num_of_Delayed_Payment	1.000000	0.315604
Changed_Credit_Limit	0.315604	1.000000
Num_Credit_Inquiries	0.497532	0.368664
Outstanding_Debt	0.505241	0.464225
Credit_Utilization_Ratio	-0.062507	-0.038766
Credit_History_Age	-0.481429	-0.425295
Total_EMI_per_month	-0.007373	0.001202
Amount_invested_monthly	-0.024517	-0.007301
Monthly_Balance	-0.294056	-0.195844

	Num_Credit_Inquiries	Outstanding_Debt \
Age	-0.243182	-0.200344
Annual_Income	-0.013630	-0.008690
Monthly_Inhand_Salary	-0.285780	-0.277441
Num_Bank_Accounts	0.526000	0.515963
Num_Credit_Card	0.468567	0.498871
Interest_Rate	0.633609	0.633570
Num_of_Loan	0.544718	0.640329
Delay_from_due_date	0.539331	0.578490
Num_of_Delayed_Payment	0.497532	0.505241
Changed_Credit_Limit	0.368664	0.464225
Num_Credit_Inquiries	1.000000	0.597826
Outstanding_Debt	0.597826	1.000000
Credit_Utilization_Ratio	-0.066738	-0.058116
Credit_History_Age	-0.603835	-0.632136
Total_EMI_per_month	-0.005862	-0.001622
Amount_invested_monthly	-0.009351	-0.012362
Monthly_Balance	-0.311252	-0.316415

	Credit_Utilization_Ratio	Credit_History_Age \
Age	0.022428	0.232788
Annual_Income	0.009059	0.003094
Monthly_Inhand_Salary	0.159732	0.283447
Num_Bank_Accounts	-0.060204	-0.487774
Num_Credit_Card	-0.048869	-0.419828
Interest_Rate	-0.066029	-0.576757
Num_of_Loan	-0.077027	-0.596080
Delay_from_due_date	-0.054367	-0.487857
Num_of_Delayed_Payment	-0.062507	-0.481429
Changed_Credit_Limit	-0.038766	-0.425295

Num_Credit_Inquiries	-0.066738	-0.603835
Outstanding_Debt	-0.058116	-0.632136
Credit_Utilization_Ratio	1.000000	0.058006
Credit_History_Age	0.058006	1.000000
Total_EMI_per_month	0.000988	0.004571
Amount_invested_monthly	0.014140	0.012146
Monthly_Balance	0.216091	0.326325

	Total_EMI_per_month	Amount_invested_monthly \
Age	0.000514	-0.004075
Annual_Income	-0.010581	-0.000604
Monthly_Inhand_Salary	0.010987	0.049266
Num_Bank_Accounts	0.005370	-0.013920
Num_Credit_Card	-0.001450	-0.001580
Interest_Rate	0.003615	-0.010806
Num_of_Loan	-0.000914	-0.019284
Delay_from_due_date	-0.005720	-0.004256
Num_of_Delayed_Payment	-0.007373	-0.024517
Changed_Credit_Limit	0.001202	-0.007301
Num_Credit_Inquiries	-0.005862	-0.009351
Outstanding_Debt	-0.001622	-0.012362
Credit_Utilization_Ratio	0.000988	0.014140
Credit_History_Age	0.004571	0.012146
Total_EMI_per_month	1.000000	0.001333
Amount_invested_monthly	0.001333	1.000000
Monthly_Balance	0.003377	-0.008427

	Monthly_Balance
Age	0.108362
Annual_Income	0.020047
Monthly_Inhand_Salary	0.680436
Num_Bank_Accounts	-0.284251
Num_Credit_Card	-0.235929
Interest_Rate	-0.318435
Num_of_Loan	-0.406558
Delay_from_due_date	-0.267302
Num_of_Delayed_Payment	-0.294056
Changed_Credit_Limit	-0.195844
Num_Credit_Inquiries	-0.311252
Outstanding_Debt	-0.316415
Credit_Utilization_Ratio	0.216091
Credit_History_Age	0.326325
Total_EMI_per_month	0.003377
Amount_invested_monthly	-0.008427
Monthly_Balance	1.000000

We can see that some of the correlations are unusually high, for example, 0.680436 (close to or

greater than 0.6) between monthly balance and monthly inhand salary. This might be causing multi-collinearity issues down the line. Some of the solutions which we can employ in this case range from removing the highly correlated independent variables to introducing a linear combination of the highly correlated independent variables instead of having them individually present as columns. I am deleting the column for monthly balance below to resolve this issue.

```
[8]: del data["Monthly_Balance"]

'''
We also delete the identifying information such as SSN, ID, Customer_ID, and
↳Name
before converting our categorical variables to numerical encoding because we
↳don't
want these columns to be encoded as this might cause overfitting and it might
↳lead
to non-convergence issues with our model because each observation will take on
↳a
different value and number of categories will be very close to the total number
↳of
observations.
'''
del data["SSN"]
del data["ID"]
del data["Customer_ID"]
del data["Name"]

data
```

```
[8]:
```

	Month	Age	Occupation	Annual_Income	Monthly_Inhand_Salary	\
0	July	23	Scientist	19114	1824.843333	
1	February	28	Teacher	34847	3037.986667	
2	May	28	Teacher	34847	3037.986667	
3	June	28	Teacher	34847	3037.986667	
4	August	28	Teacher	34847	3037.986667	
...	
26523	January	18	Lawyer	42903	3468.315833	
26524	January	49	Writer	37188	3097.008333	
26525	February	49	Writer	37188	3097.008333	
26526	February	28	Architect	20002	1929.906667	
26527	August	29	Architect	20002	1929.906667	

	Num_Bank_Accounts	Num_Credit_Card	Interest_Rate	Num_of_Loan	\
0	3	4	3	4	
1	2	4	6	1	
2	2	4	6	1	
3	2	4	6	1	
4	2	4	6	1	

...
26523	0	4	6	1
26524	1	4	5	3
26525	1	4	5	3
26526	10	8	29	5
26527	10	8	29	5

	Type_of_Loan	...	\
0	Auto Loan, Credit-Builder Loan, Personal Loan,...	...	
1	Credit-Builder Loan	...	
2	Credit-Builder Loan	...	
3	Credit-Builder Loan	...	
4	Credit-Builder Loan	...	
...	
26523	Not Specified	...	
26524	Home Equity Loan, Mortgage Loan, and Student Loan	...	
26525	Home Equity Loan, Mortgage Loan, and Student Loan	...	
26526	Personal Loan, Auto Loan, Mortgage Loan, Stude...	...	
26527	Personal Loan, Auto Loan, Mortgage Loan, Stude...	...	

	Num_Credit_Inquiries	Credit_Mix	Outstanding_Debt	\
0	4.0	Good	809.98	
1	2.0	Good	605.03	
2	2.0	Good	605.03	
3	2.0	Good	605.03	
4	2.0	Good	605.03	
...	
26523	1.0	Good	1079.48	
26524	3.0	Good	620.64	
26525	3.0	Good	620.64	
26526	9.0	Bad	3571.70	
26527	9.0	Bad	3571.70	

	Credit_Utilization_Ratio	Credit_History_Age	Payment_of_Min_Amount	\
0	22.537593	22.583333	No	
1	38.550848	26.666667	No	
2	34.977895	26.916667	No	
3	33.381010	27.000000	No	
4	32.933856	27.166667	No	
...	
26523	27.289440	28.083333	No	
26524	39.080823	29.750000	No	
26525	32.803431	29.833333	No	
26526	39.772607	5.750000	Yes	
26527	37.140784	6.250000	Yes	

Total_EMI_per_month	Amount_invested_monthly	\
---------------------	-------------------------	---

0	49.574949	178.344067
1	18.816215	40.391238
2	18.816215	130.115420
3	18.816215	43.477190
4	18.816215	218.904344
...
26523	50894.000000	78.514945
26524	84.205949	223.875018
26525	84.205949	70.869970
26526	12112.000000	148.275233
26527	60.964772	34.662906

	Payment_Behaviour	Credit_Score
0	Low_spent_Small_value_payments	Good
1	High_spent_Large_value_payments	Good
2	Low_spent_Small_value_payments	Good
3	High_spent_Large_value_payments	Good
4	Low_spent_Small_value_payments	Good
...
26523	High_spent_Small_value_payments	Good
26524	Low_spent_Small_value_payments	Good
26525	High_spent_Large_value_payments	Good
26526	Low_spent_Small_value_payments	Poor
26527	High_spent_Large_value_payments	Standard

[26528 rows x 23 columns]

- Appropriately encode all categorical features in the dataframe

```
[9]: data.dtypes
```

```
[9]: Month                object
     Age                  int64
     Occupation           object
     Annual_Income        int64
     Monthly_Inhand_Salary float64
     Num_Bank_Accounts    int64
     Num_Credit_Card      int64
     Interest_Rate        int64
     Num_of_Loan          int64
     Type_of_Loan         object
     Delay_from_due_date  int64
     Num_of_Delayed_Payment int64
     Changed_Credit_Limit float64
     Num_Credit_Inquiries float64
     Credit_Mix           object
     Outstanding_Debt     float64
```

```

Credit_Utilization_Ratio    float64
Credit_History_Age          float64
Payment_of_Min_Amount       object
Total_EMI_per_month         float64
Amount_invested_monthly     float64
Payment_Behaviour           object
Credit_Score                object
dtype: object

```

```

[10]: from sklearn.preprocessing import LabelEncoder

# get the categorical columns
cat_cols = data.select_dtypes(include=['object']).columns

# create the label encoder object
label_encoder = LabelEncoder()

# apply label encoding to each categorical column
data[cat_cols] = data[cat_cols].apply(lambda col: label_encoder.
    ↪fit_transform(col))

# print the encoded dataframe
data

```

```

[10]:
      Month  Age  Occupation  Annual_Income  Monthly_Inhand_Salary  \
0         4   23          12         19114         1824.843333
1         2   28          13         34847         3037.986667
2         7   28          13         34847         3037.986667
3         5   28          13         34847         3037.986667
4         1   28          13         34847         3037.986667
...      ...  ...      ...      ...      ...
26523     3   18           7         42903         3468.315833
26524     3   49          14         37188         3097.008333
26525     2   49          14         37188         3097.008333
26526     2   28           1         20002         1929.906667
26527     1   29           1         20002         1929.906667

      Num_Bank_Accounts  Num_Credit_Card  Interest_Rate  Num_of_Loan  \
0                    3                4                3            4
1                    2                4                6            1
2                    2                4                6            1
3                    2                4                6            1
4                    2                4                6            1
...                  ...              ...              ...          ...
26523                 0                4                6            1
26524                 1                4                5            3
26525                 1                4                5            3

```

26526	10	8	29	5
26527	10	8	29	5

	Type_of_Loan	...	Num_Credit_Inquiries	Credit_Mix	Outstanding_Debt	\
0	119	...	4.0	1	809.98	
1	649	...	2.0	1	605.03	
2	649	...	2.0	1	605.03	
3	649	...	2.0	1	605.03	
4	649	...	2.0	1	605.03	
...	
26523	3257	...	1.0	1	1079.48	
26524	2335	...	3.0	1	620.64	
26525	2335	...	3.0	1	620.64	
26526	4636	...	9.0	0	3571.70	
26527	4636	...	9.0	0	3571.70	

	Credit_Utilization_Ratio	Credit_History_Age	Payment_of_Min_Amount	\
0	22.537593	22.583333	0	
1	38.550848	26.666667	0	
2	34.977895	26.916667	0	
3	33.381010	27.000000	0	
4	32.933856	27.166667	0	
...	
26523	27.289440	28.083333	0	
26524	39.080823	29.750000	0	
26525	32.803431	29.833333	0	
26526	39.772607	5.750000	1	
26527	37.140784	6.250000	1	

	Total_EMI_per_month	Amount_invested_monthly	Payment_Behaviour	\
0	49.574949	178.344067	5	
1	18.816215	40.391238	0	
2	18.816215	130.115420	5	
3	18.816215	43.477190	0	
4	18.816215	218.904344	5	
...	
26523	50894.000000	78.514945	2	
26524	84.205949	223.875018	5	
26525	84.205949	70.869970	0	
26526	12112.000000	148.275233	5	
26527	60.964772	34.662906	0	

	Credit_Score
0	0
1	0
2	0
3	0

```

4          0
...
26523      0
26524      0
26525      0
26526      1
26527      2

```

```
[26528 rows x 23 columns]
```

```
[11]: data.dtypes
```

```

[11]: Month          int64
Age                int64
Occupation         int64
Annual_Income      int64
Monthly_Inhand_Salary  float64
Num_Bank_Accounts  int64
Num_Credit_Card    int64
Interest_Rate      int64
Num_of_Loan        int64
Type_of_Loan       int64
Delay_from_due_date int64
Num_of_Delayed_Payment int64
Changed_Credit_Limit float64
Num_Credit_Inquiries float64
Credit_Mix         int64
Outstanding_Debt   float64
Credit_Utilization_Ratio float64
Credit_History_Age float64
Payment_of_Min_Amount int64
Total_EMI_per_month float64
Amount_invested_monthly float64
Payment_Behaviour  int64
Credit_Score       int64
dtype: object

```

3. Logistic Regression (20 points)

- **What is the use of validation dataset in machine learning?** The validation dataset is a subset of the original dataset that is used to evaluate the performance of a model during training and to select the best hyperparameters.

During the training phase of a machine learning model, the model is trained on a training dataset, and its performance is evaluated on a validation dataset. The goal is to optimize the model's performance on the validation dataset while preventing overfitting to the training dataset. Overfitting occurs when the model learns the noise or the specifics of the training dataset too well, resulting

in poor generalization to new, unseen data.

The validation dataset is used to tune the hyperparameters of the model, such as the learning rate, the number of layers, the number of hidden units, etc. Hyperparameters are model parameters that are not learned during training and must be set before training the model. By evaluating the model's performance on the validation dataset for different hyperparameter values, we can select the best hyperparameters that optimize the model's performance.

The main advantage of using a validation dataset is that it provides an unbiased estimate of the model's performance on unseen data. This estimate can be used to compare different models or to select the best hyperparameters for a given model. Without a validation dataset, it is difficult to estimate the model's performance on unseen data accurately, and we risk overfitting the model to the training data.

To summarize, the validation dataset is used in machine learning to:

- Evaluate the performance of the model during training
 - Prevent overfitting of the model to the training data
 - Tune the hyperparameters of the model
 - Provide an unbiased estimate of the model's performance on unseen data.
- **Make an appropriate split of the data train, validation, test and fit a multi-class logistic regression model using the Scikit-learn library.**

```
[22]: from sklearn.model_selection import train_test_split
      from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import accuracy_score
      import pandas as pd

      # Split the data into training and test sets
      X_train, X_test, y_train, y_test = train_test_split(data.drop('Credit_Score',
      ↪axis=1), data['Credit_Score'], test_size=0.2, random_state=8990)

      # Split the training data into training and validation sets
      X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.
      ↪2, random_state=8990)

      # Fit the multi-class logistic regression model
      clf = LogisticRegression(multi_class='multinomial', solver='lbfgs',
      ↪max_iter=10000)
      clf.fit(X_train, y_train)

      # Predict the classes of the validation set
      y_val_pred = clf.predict(X_val)

      # Calculate the validation set accuracy
      val_accuracy = accuracy_score(y_val, y_val_pred)

      # Predict the classes of the test set
      y_test_pred = clf.predict(X_test)
```



```
# Calculate the test set accuracy
test_accuracy = accuracy_score(y_test, y_test_pred)

print(f"Validation set accuracy: {val_accuracy:.3f}")
print(f"Test set accuracy: {test_accuracy:.3f}")
```

Validation set accuracy: 0.554

Test set accuracy: 0.543

- **Print the classification report of the model. What do you find in the report?**

```
[23]: from sklearn.metrics import classification_report

logistic_model_classification_report = classification_report(y_test,
    ↪ y_test_pred)

print(logistic_model_classification_report)
```

	precision	recall	f1-score	support
0	0.35	0.01	0.02	854
1	0.51	0.35	0.42	1626
2	0.55	0.81	0.66	2826
accuracy			0.54	5306
macro avg	0.47	0.39	0.36	5306
weighted avg	0.51	0.54	0.48	5306

This classification report is for a logistic regression model, which has been used to classify data into three classes (0, 1, and 2).

The report shows various metrics that evaluate the performance of the model. Here is a brief explanation of each metric:

- **Precision:** Precision measures the proportion of true positives (correctly predicted instances of a class) among the instances that the model predicted as positive. In this report, the precision for class 0 is very low at 0.35, indicating that the model has a high number of false positives for this class. The precision for class 1 and 2 are 0.51 and 0.55, respectively, indicating that the model is better at predicting these classes.
- **Recall:** Recall measures the proportion of true positives among the instances that actually belong to a class. In this report, the recall for class 0 is very low at 0.01, indicating that the model is not good at identifying instances of this class. The recall for class 1 and 2 are 0.35 and 0.81, respectively, indicating that the model is better at identifying these classes.
- **F1-score:** F1-score is a weighted average of precision and recall that takes into account both metrics. It is a good metric for evaluating the overall performance of a model. In this report, the F1-score for class 0 is very low at 0.02, indicating that the model is not good at predicting

this class. The F1-scores for class 1 and 2 are 0.42 and 0.66, respectively, indicating that the model is better at predicting these classes.

- **Support:** Support indicates the number of instances of each class in the test set.
- **Accuracy:** Accuracy measures the proportion of correctly classified instances among all instances. In this report, the overall accuracy of the model is 0.54, indicating that the model is correct in its prediction for 54% of instances.
- **Macro average:** Macro average calculates the average metric score across all classes, giving equal weight to each class. In this report, the macro average F1-score is 0.36, indicating that the overall performance of the model is not very good.
- **Weighted average:** Weighted average calculates the average metric score across all classes, weighting each class by its support. In this report, the weighted average F1-score is 0.48, which is slightly better than the macro average F1-score but still indicates that the model's performance is not very good.

Overall, this classification report suggests that the logistic regression model is not very accurate at predicting class 0, and performs better at predicting classes 1 and 2. The overall performance of the model is also not very good, with a low macro average F1-score and accuracy of 54%.

4 4. Decision Tree and Bagging (30 points)

• **Describe the algorithm CART. What are the advantages and disadvantages of CART?** CART (Classification And Regression Trees) is a decision tree algorithm used for both classification and regression tasks. The algorithm works by recursively splitting the data into two subsets, based on a single feature at a time, until the subsets are as homogeneous as possible in terms of the target variable.

Here are the main steps of the CART algorithm:

1. Select the feature that provides the best split. The feature with the highest information gain (for classification) or the highest reduction in variance (for regression) is chosen as the splitting feature.
2. Split the data into two subsets based on the chosen feature. The subsets are chosen such that each subset is as pure as possible with respect to the target variable.
3. Repeat steps 1 and 2 for each subset, until a stopping criterion is met. The stopping criterion could be a maximum depth of the tree, a minimum number of samples per leaf, or a minimum improvement in impurity.
4. Create a tree by assigning a label or value to each leaf node, based on the majority class (for classification) or the mean value (for regression) of the samples in the leaf node.

Advantages of CART:

1. CART is a simple and interpretable algorithm. The resulting decision tree can be easily visualized and understood, which is useful for explaining the reasoning behind the model's predictions.
2. CART can handle a mix of categorical and continuous features, making it versatile and applicable to a wide range of problems.

3. CART is a non-parametric algorithm, which means that it does not make any assumptions about the distribution of the data. This makes it useful when the data does not follow a specific distribution.
4. CART can handle missing values and outliers by using surrogate splits, which improve the robustness of the model.

Disadvantages of CART:

1. CART is prone to overfitting the data, which means that it can learn the noise or specific features of the training data too well, resulting in poor generalization to new, unseen data.
2. CART is sensitive to small variations in the data and can produce different trees for different splits, which makes the model less stable and harder to interpret.
3. CART is a greedy algorithm, which means that it makes the best split at each step without considering the overall structure of the tree. This can lead to suboptimal solutions.
4. CART can be biased towards features with many categories, as they can dominate the splitting process and result in an unbalanced tree.

Overall, CART is a powerful algorithm that can be used for a wide range of problems. However, to avoid overfitting and improve the performance of the model, it is important to tune the hyperparameters of the algorithm and use techniques such as pruning, regularization, and ensemble methods.

- **Implement DecisionTreeClassifier from the sklearn library to train one decision tree. You can evaluate the accuracy of the validation set to tune model parameters. You should only evaluate your final accuracy on the test dataset.**

```
[24]: from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

# Split data into train, validation, and test sets
train_data, test_data, train_labels, test_labels = train_test_split(data,
    ↳ drop('Credit_Score', axis=1), data['Credit_Score'], test_size=0.2,
    ↳ random_state=8990)
train_data, val_data, train_labels, val_labels = train_test_split(train_data,
    ↳ train_labels, test_size=0.2, random_state=8990)

# Train a decision tree classifier on the training data
clf = DecisionTreeClassifier()
clf.fit(train_data, train_labels)

# Evaluate the accuracy on the validation set to tune model parameters
val_pred = clf.predict(val_data)
val_acc = accuracy_score(val_labels, val_pred)
print("Validation accuracy:", val_acc)

# Evaluate the accuracy on the test set
```

```
test_pred = clf.predict(test_data)
test_acc = accuracy_score(test_labels, test_pred)
print("Test accuracy:", test_acc)
```


Validation accuracy: 0.6525323910482921

Test accuracy: 0.6515265736901621

- Use sklearn's `sklearn.tree.plot_tree` method and matplotlib to visualize your classification tree.

```
[25]: from sklearn.tree import plot_tree
import matplotlib.pyplot as plt

# Visualize the decision tree
plt.figure()
plot_tree(clf, filled=True)
plt.show()
```



- Use 30 different random seeds to train 30 identical decision trees and record the test accuracies. Calculate and report the average accuracy and standard deviation across the 30 runs. What do you find using this bagging method?

```
[26]: import numpy as np
```

```

# Define a function to train a decision tree using a given random seed
def train_decision_tree(seed, train_data, train_labels, test_data, test_labels):
    clf = DecisionTreeClassifier(random_state=seed)
    clf.fit(train_data, train_labels)
    test_pred = clf.predict(test_data)
    test_acc = accuracy_score(test_labels, test_pred)
    return test_acc

# Split data into train and test sets
train_data, test_data, train_labels, test_labels = train_test_split(data,
    ↳drop('Credit_Score', axis=1), data['Credit_Score'], test_size=0.2,
    ↳random_state=8990)

# Train 30 decision trees with different random seeds and record the test
    ↳accuracies
test_accs = []
for seed in range(30):
    test_acc = train_decision_tree(seed, train_data, train_labels, test_data,
    ↳test_labels)
    test_accs.append(test_acc)

# Calculate and report the average accuracy and standard deviation across the
    ↳30 runs
avg_acc = np.mean(test_accs)
std_dev = np.std(test_accs)
print("Average accuracy:", avg_acc)
print("Standard deviation:", std_dev)

```

Average accuracy: 0.6578967206935545
Standard deviation: 0.0030367385530887395

Bagging, or bootstrap aggregation, is a technique that involves training multiple models on different subsets of the training data and combining their predictions to obtain a final prediction. In this case, we are training identical decision trees on different subsets of the data by using different random seeds, and averaging their test accuracies to obtain an estimate of the overall accuracy of the model.

By using this bagging method, we can obtain a more robust estimate of the accuracy of the model, and reduce the impact of random variations in the data or the training process. The standard deviation across the 30 runs can give us an idea of the level of variability in the accuracy estimates, and help us assess the reliability of the model.

In this case specifically, we found that the standard deviation is very low and the average accuracy is close to the accuracy we obtained in the first run in the part before this.

5 5. Random Forest (25 points)

• **What is the difference between bagging and random forest?** Bagging and random forests are both ensemble learning techniques that combine the predictions of multiple decision trees to improve the accuracy and robustness of the model. However, there are some differences between these two methods:

- **Sampling method:** Bagging and random forests use different sampling methods to generate the training sets for the individual decision trees. Bagging samples the training data with replacement to create multiple subsets, each of which is used to train a decision tree. Random forests also sample the features used to split the nodes of each decision tree, in addition to sampling the training data.
- **Decision tree construction:** Bagging and random forests use different techniques for constructing the individual decision trees. Bagging uses a standard decision tree algorithm, such as CART or ID3, to build each tree independently. Random forests use a modified version of the decision tree algorithm, where only a random subset of features is considered for each split, rather than all features.
- **Prediction aggregation:** Bagging and random forests use different methods for aggregating the predictions of the individual decision trees. Bagging usually takes a simple average of the predictions of all decision trees to generate the final prediction. Random forests use a majority vote or weighted average of the predictions, where the weight of each tree depends on its accuracy.
- **Performance:** Random forests tend to perform better than bagging, especially for high-dimensional datasets with many features, because they can reduce the variance and overfitting of the individual decision trees. Random forests also provide an estimate of the importance of each feature, which can be useful for feature selection and interpretation.

In summary, both bagging and random forests are powerful ensemble learning techniques that can improve the accuracy and robustness of decision tree models. However, random forests add an additional layer of randomness by sampling the features used to split the nodes, which can further reduce the variance and overfitting of the individual trees, and provide better performance for high-dimensional datasets.

• **Why is it important for individual estimators in the random forest to have access to only a subset of all features?** In random forests, it is important for individual estimators (i.e., decision trees) to have access to only a subset of all features for several reasons:

- **Reduces correlation:** When decision trees are trained on all features, they tend to be highly correlated and make similar predictions, which can lead to overfitting and reduced performance. By limiting the number of features available to each tree, random forests can reduce the correlation between the trees and improve the diversity of the ensemble.
- **Reduces overfitting:** When there are many features in the dataset, individual trees can easily overfit the training data and perform poorly on unseen data. By using a random subset of features for each tree, random forests can reduce the overfitting and improve the generalization performance of the ensemble.
- **Faster training:** When there are many features in the dataset, training decision trees on all features can be computationally expensive and time-consuming. By using a random subset

of features for each tree, random forests can reduce the training time and make the algorithm more scalable.

- **Feature importance:** Random forests can also provide an estimate of the importance of each feature in the dataset, based on how often they are used to split the nodes of the individual trees. By identifying the most important features, random forests can help with feature selection and interpretation, and improve the understanding of the underlying data.

In summary, limiting the number of features available to each decision tree in random forests can improve the diversity, generalization performance, training time, and interpretability of the ensemble, making it a powerful and widely used machine learning technique.

- **Implement a random forest classifier to solve the classification problem again.**

```
[27]: from sklearn.ensemble import RandomForestClassifier
      from sklearn.model_selection import train_test_split
      from sklearn.metrics import accuracy_score

      # Split the data into training and testing sets
      X_train, X_test, y_train, y_test = train_test_split(data.drop('Credit_Score',
      ↪axis=1), data['Credit_Score'], test_size=0.2, random_state=8990)

      # Initialize the random forest classifier
      rfc = RandomForestClassifier(n_estimators=100, max_features='sqrt',
      ↪random_state=8990)

      # Fit the classifier to the training data
      rfc.fit(X_train, y_train)

      # Make predictions on the test data
      y_pred = rfc.predict(X_test)

      # Evaluate the accuracy of the model
      accuracy = accuracy_score(y_test, y_pred)
      print("Accuracy:", accuracy)
```

Accuracy: 0.7664907651715039

• **Compare and contrast models from Section 3, 4 and 5.** Logistic regression, decision tree, and random forest are all popular machine learning algorithms used for classification tasks. Here are some key differences between these models:

- **Model complexity:** Logistic regression is a linear model that tries to find a linear boundary between classes, while decision trees and random forests can model more complex nonlinear boundaries. Random forests are a type of ensemble learning method that combines multiple decision trees, which can improve the accuracy and reduce overfitting compared to a single decision tree.
- **Interpretability:** Logistic regression is a simple and interpretable model that can provide coefficients that indicate the importance of each feature. Decision trees can also provide in-

sights into the decision-making process by showing the split points for each feature. However, random forests can be more difficult to interpret due to the large number of trees and the ensemble nature of the model.

- ***Robustness to noise:*** Decision trees are susceptible to overfitting and can be sensitive to small changes in the training data. Random forests can reduce overfitting and improve generalization performance by using a subset of features and a bootstrap sampling technique. Logistic regression can also be susceptible to overfitting, but can be more robust to noisy features.
- ***Scalability:*** Logistic regression is a relatively fast and scalable algorithm, especially when the number of features is small. Decision trees and random forests can be slower and more computationally intensive, especially when the number of features and trees is large.

In terms of performance, the best algorithm depends on the specific dataset and problem at hand. Logistic regression can be effective when the relationship between the features and target is linear, while decision trees and random forests can handle more complex relationships. Random forests can often provide the highest accuracy and are robust to overfitting, but can be more difficult to interpret and slower to train.

In summary, logistic regression, decision tree, and random forest are all useful algorithms for classification tasks, but differ in terms of model complexity, interpretability, robustness to noise, and scalability. It is important to experiment with different algorithms and evaluate their performance on the specific dataset to determine the best model for the task at hand.

In this specific case, we can see that the Random Forests model gives us the highest accuracy of around 0.76. We also note that the accuracy of random forest model is almost 0.2 greater than that of logistic regression and almost 0.10 more than that obtained with the decision trees and bagging method.