

GR5260

# **Programming for Quantitative & Computational Finance**

Instructor: Ka Yi Ng

April 21, 2023

## **Tree-based models**

# Tree-based ML models


- Partitioning feature space into rectangular blocks, assigning a constant prediction value to each block
- Decision trees using CART algorithm
- Random forests: ensemble example
- Recall:
  - Training dataset:  $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)}), \dots, (\mathbf{x}^{(m)}, \mathbf{y}^{(m)})$
  - Features:  $\mathbf{x} = (x_1, \dots, x_n)$
  - Classification:  $\mathbf{y} \in \{1, 2, \dots, K\}$
  - Regression:  $\mathbf{y} \in \mathbb{R}$

# Classification: CART algorithm

- Gini impurity index for a set of  $N$  points in  $K$  classes:

- $$G = \sum_{i=1}^K \hat{p}_i(1 - \hat{p}_i), \quad \hat{p}_i = \frac{N_{class\_i}}{N} \text{ and } N_{class\_i} = \# \text{ points of class } i$$

- As an error measure  $\hat{p}_{blue}(1 - \hat{p}_{blue})$



- 3 red points, 3 blue points, 7 green points


- $G = \left(\frac{3}{13}\right)\left(\frac{10}{13}\right) + \left(\frac{3}{13}\right)\left(\frac{10}{13}\right) + \left(\frac{7}{13}\right)\left(\frac{6}{13}\right) = 0.60$

- 3 red points, 4 blue points, 6 green points

- $G = \left(\frac{3}{13}\right)\left(\frac{10}{13}\right) + \left(\frac{4}{13}\right)\left(\frac{9}{13}\right) + \left(\frac{6}{13}\right)\left(\frac{7}{13}\right) = 0.64$

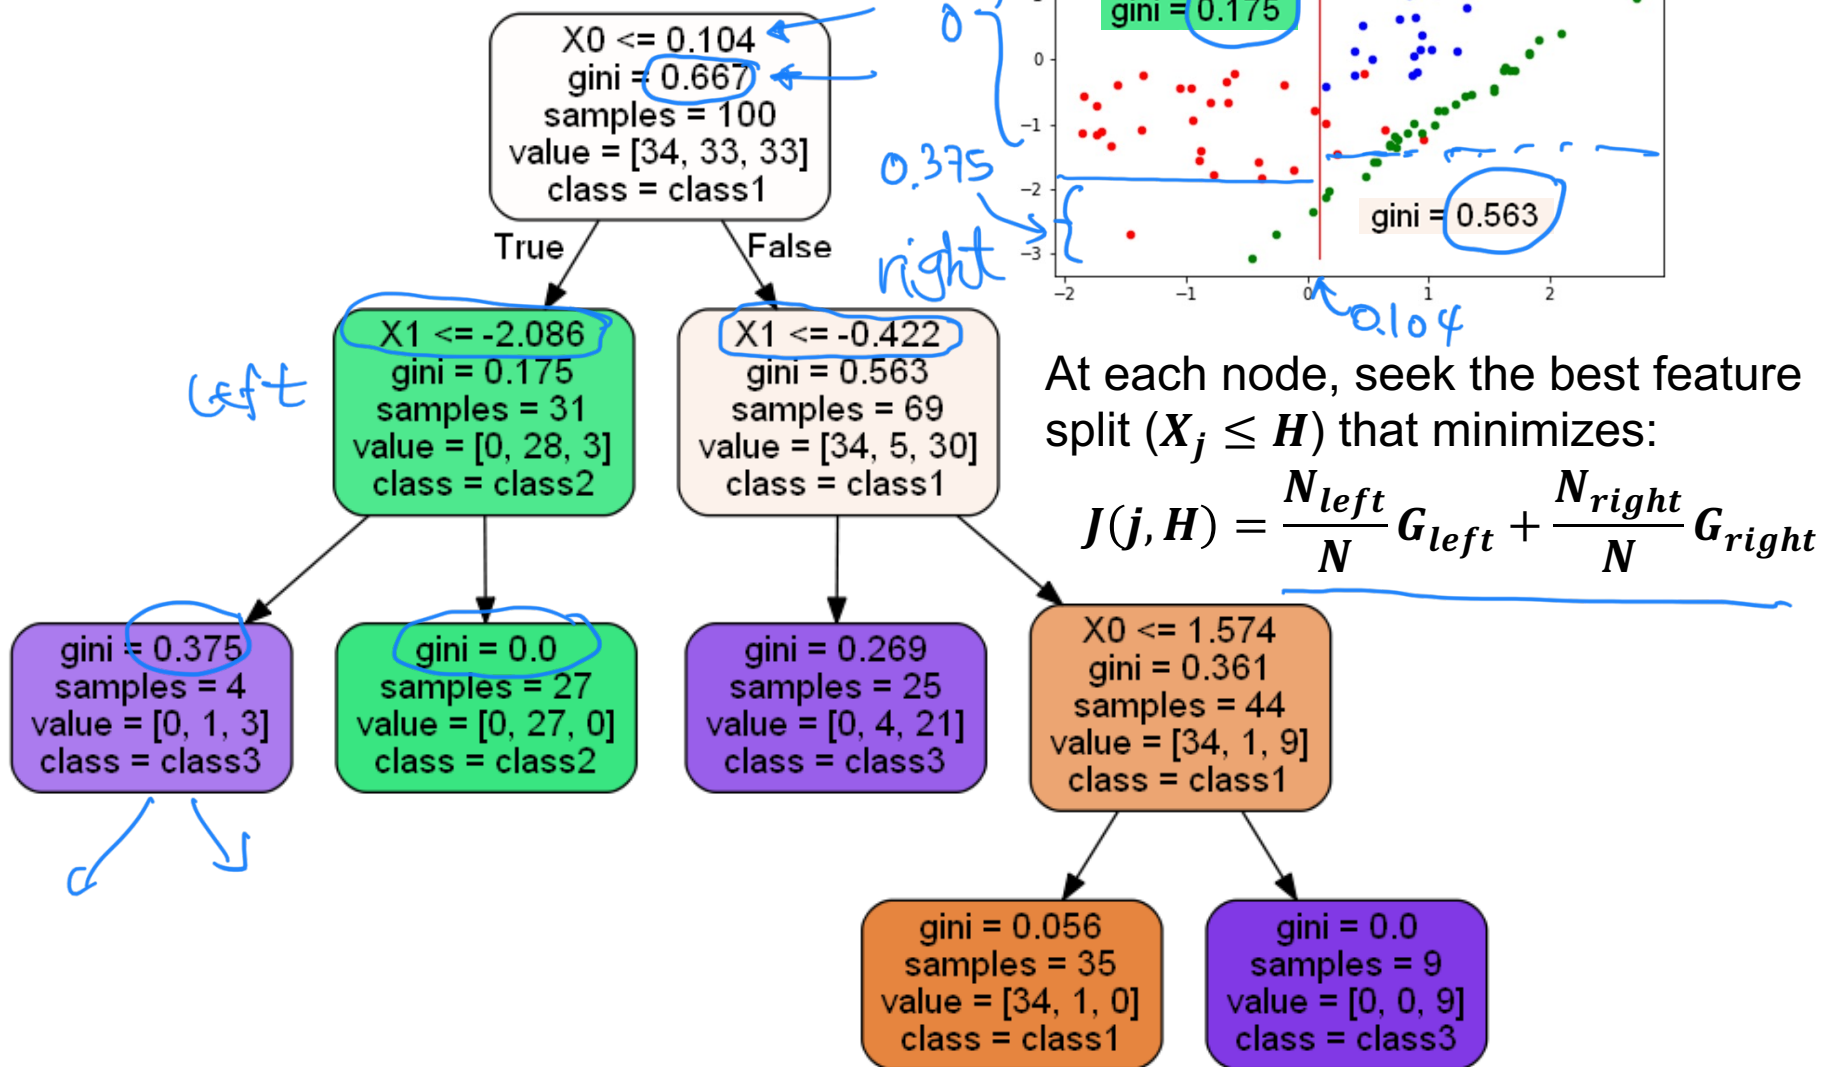
# Decision Tree: Classification

Partition n-dim feature space using hyperplanes defined by feature values



# Classification: CART algorithm

- Gini impurity index  $G = \sum_{i=1}^K \hat{p}_i(1 - \hat{p}_i)$



# Classification: CART algorithm

- At each node:
    - Compute its Gini impurity index,  $G_{node}$
    - Among all features, find the feature ( $X_j$ ) and threshold level  $H$  that gives the best split ( $X_j \leq H$ ) ie. minimizes the cost function
- $$J(j, H) = \frac{N_{left}}{N} G_{left} + \frac{N_{right}}{N} G_{right}$$
- If the difference  $G_{node} - J(\hat{j}, \hat{H}) > \delta$ , split the node using the best split criterion; otherwise, no split
- Repeat the above until no nodes can be split or some predefined criterion (eg. #nodes, tree depth) is met
  - Prediction:
    - For each leaf node: predicted class probability distribution
    - Given a new data point  $x_{new}$ , determine the leaf node it belongs to.
    - Predicted label = class with highest probability in the leaf node

# Regression: CART algorithm

- At each node:
  - $\bar{y}_{node}$  = mean of label values of all points belonging to the node

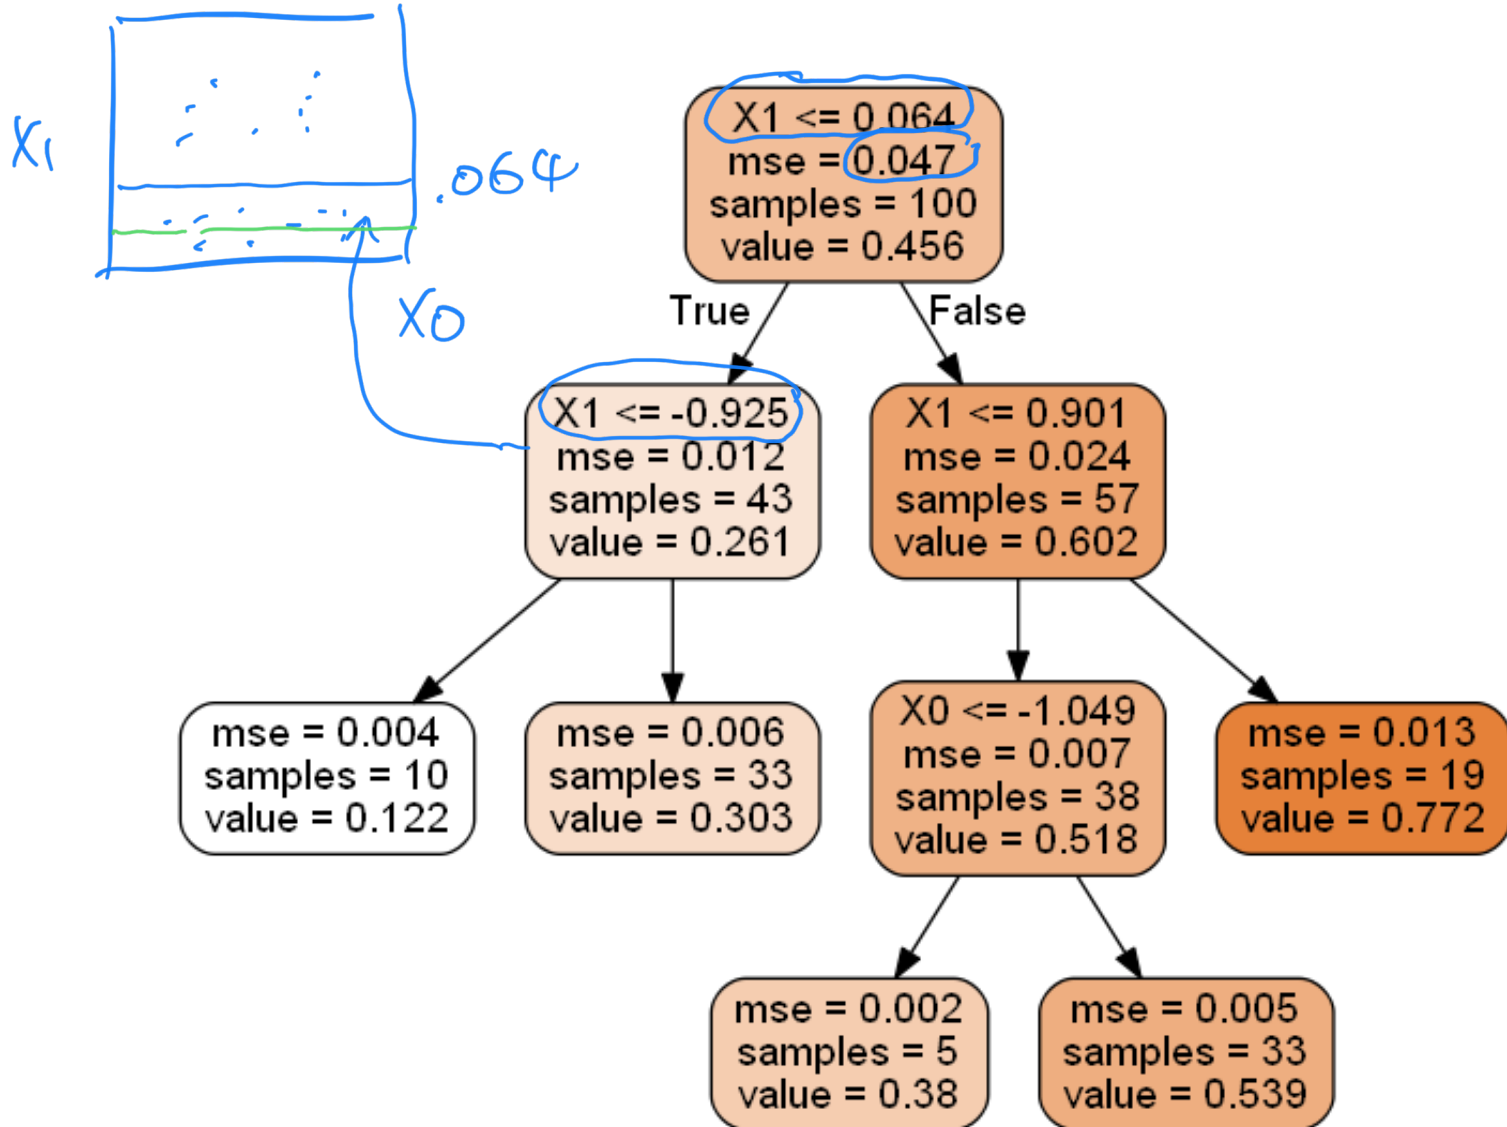
$$MSE_{node} = \frac{1}{N} \sum_{i \in node} (y^{(i)} - \bar{y}_{node})^2$$

- Among all features, find the feature ( $X_j$ ) and threshold level  $H$  that gives the best split ( $X_j \leq H$ ) ie. minimizes the cost function

$$J(j, H) = \frac{N_{left}}{N} MSE_{left} + \frac{N_{right}}{N} MSE_{right}$$

- If the difference  $MSE_{node} - J(\hat{j}, \hat{H}) > \delta$ , split the node using the best split criterion; otherwise, no split
- Repeat the above until no nodes can be split or some predefined criterion (eg. #nodes, tree depth) is met
- Prediction:
  - Given a new data point  $x_{new}$ , find the leaf node  $\mathcal{N}$  it belongs to.
  - Predicted  $y_{new}$  = average label values of training points in node  $\mathcal{N}$

# Regression: CART algorithm

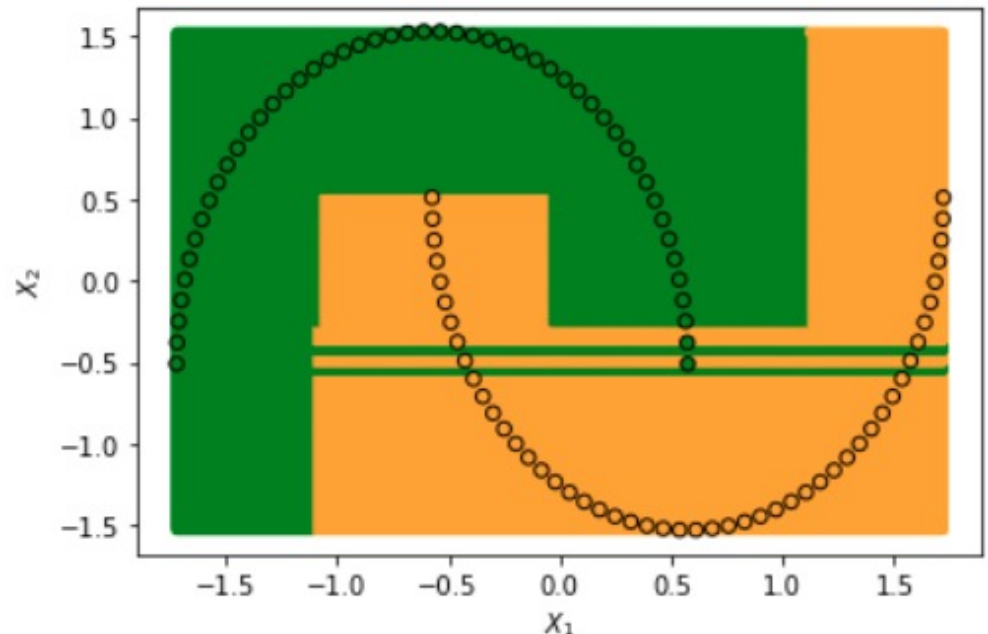





# Benefits & Limitations



- Benefits:
  - Interpretability: explicit descriptive rules
  - Flexible, no model function assumptions
- Limitations:
  - Prone to overfitting

**Moon dataset**



# Benefits & Limitations

- Limitations:
  - Sensitive to orientation of training data



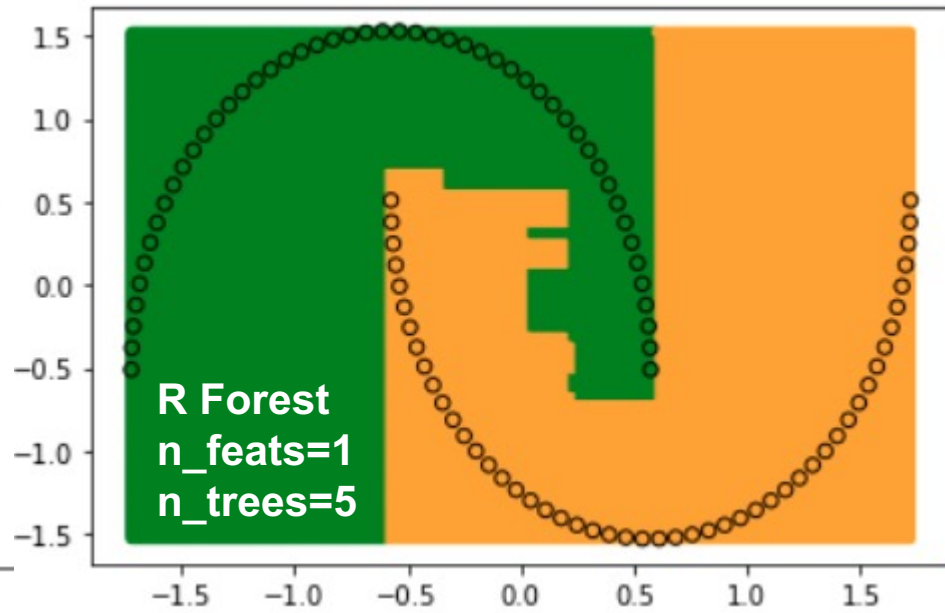
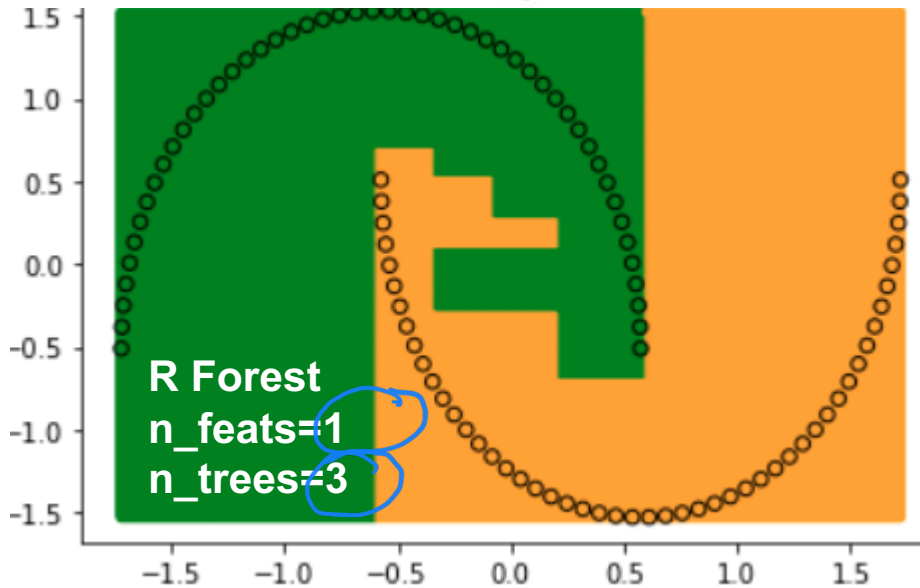

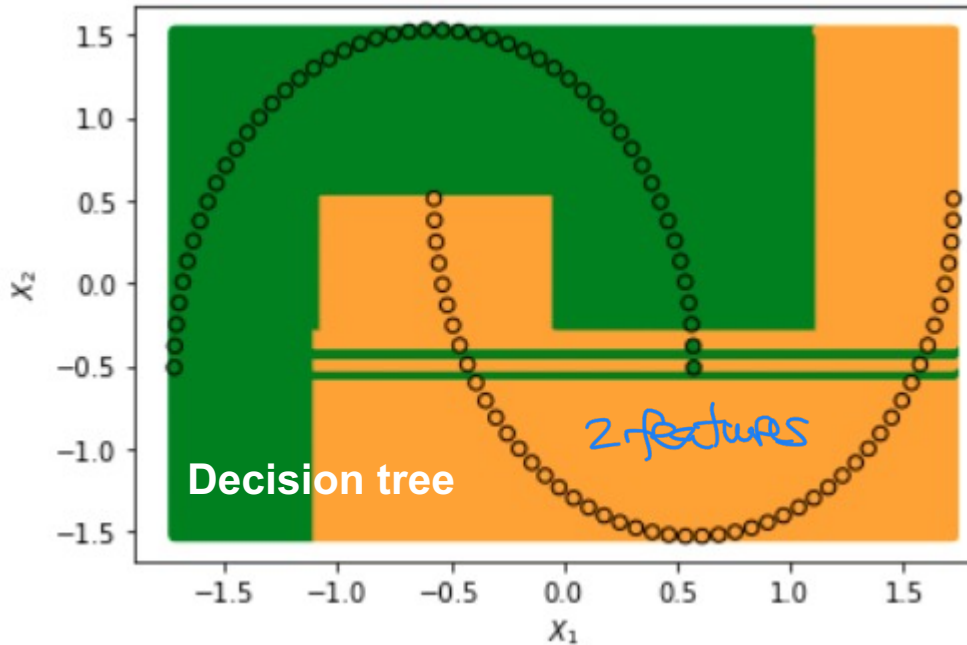
# Benefits & Limitations

- Hyperparameters: Limit size of tree to reduce overfitting
  - Tree depth
  - Maximum no. of leaf nodes
  - Minimum no. of samples in each split
  - Minimum no. of samples in each leaf node
- Random forest: reduce model instability
  - averaging prediction over multiple decision trees
- PCA: identify the principal components

# Random forests

- Uses Bagging (Bootstrap aggregation)
  - Average the predictions from a set of decision trees
  - Reduce variance (more stable prediction)
- For each tree:  $T_i$   $D_i$   $T_1, \dots, T_{500} \leftarrow$  less correlated
  - Training set: randomly sample  $m$  examples with replacement from the original training set  $D$
  - At each node: select  $\tilde{n}$  features at random ( $\tilde{n} \leq n$ ), seek best split on these features  
3 10
- Prediction: stable prediction
  - Regression: average of predicted values over the trees
  - Classification: class with highest average predicted probability over the trees

# Decision tree vs Random forest



# Benefits & Limitations

- Reduce variance and prediction error
  - Training sets: bootstrap sampling
  - Feature bagging: reduce correlation among trees
  - Average prediction
- Feature importance
  - MDI: mean decrease in impurity due to a feature
  - Alternative measures: *← applicable to any model*
  - Permutation importance (MDA): mean decrease in model performance if a feature's values on testing set are permuted

# Ensemble methods

- Train a set of models (*base learners*)
- Prediction: from combining the predictions
- Bagging
  - For base learners with high variance, low bias
  - Average prediction of models trained independently
- Boosting
  - For base learners with low variance, high bias
  - train a sequence of models, one being dependent upon the previous models
  - eg. gradient boosting: train on the residual errors made by the previous model predictor
  - Reduce high bias of weak learners; hard to scale up