



COLUMBIA UNIVERSITY
IN THE CITY OF NEW YORK

STAT 4224/5224

Bayesian Statistics

Dobrin Marchev

Recall: Metropolis-Hastings Algorithm

- Very general method for constructing a Markov chain with a specific invariant density $f(x)$, which is the target for the simulation.
- Suppose that $q(y|x)$ is any conditional density, called proposal density.
- Algorithm:

- Step 0: Start at any $X = x_0$
- Step 1: Generate $Y_i \sim q(y|x_{i-1})$
- Step 2: Set $X_i = \begin{cases} Y_i & \text{with probability } \alpha(x_{i-1}, y_i) \\ x_{i-1}, & \text{o/w} \end{cases}$

where $\alpha(x, y) = \min \left\{ \frac{f(y)}{f(x)} \frac{q(x|y)}{q(y|x)}, 1 \right\}$ is called the acceptance probability

- Go to Step 1 and repeat


Notes:

- When $q(y|x)$ is symmetric, the algorithm is known as Metropolis.
- When $q(y|x)$ does not depend on x , the algorithm is known as Independent Metropolis-Hastings.
- When $q(y|x) = g(y - x)$, the algorithm is known as Random Walk Metropolis-Hastings (see next slide)

Hamiltonian Monte Carlo

- Hamiltonian Monte Carlo improves the efficiency of MH by employing a guided proposal generation scheme.
- More specifically, HMC uses the gradient of the log posterior to direct the Markov chain towards regions of higher posterior density.
- As a result, a well-tuned HMC chain will accept proposals at a much higher rate than the traditional MH algorithm.
- While both MH and HMC produce ergodic Markov chains, the mathematics of HMC is substantially more complex than that of MH.

Main Idea



- Suppose $f(\theta)$ is a one-dimensional posterior density function, and $-\log[f(\theta)]$ assumes the shape of an inverse bell-shaped curve.
- To generate θ in a region of high posterior density, one needs to sample θ in the region corresponding to the lower values of $-\log f(\theta)$; the region can be reached with the guidance of the gradient of $-\log f(\theta)$.

Hamiltonian Dynamics



- In a sense, the approach is analogous to the movement of a hypothetical object on a frictionless curve, where the object traverses and lingers at the bottom of the valley while occasionally visiting the higher grounds on both sides.
- In classical mechanics, such movements are described by the Hamiltonian equations, where the exchanges of kinetic and potential energy dictate the object's location at any given moment.

The Hamiltonian Equations

In a Hamiltonian system, the horizontal and vertical positions are given by (θ, \mathbf{p}) . In MCMC, we are interested in $\theta \sim f(\theta)$. The parameter p , which is often referred to as the momentum, is an auxiliary quantity that we use to simulate θ under the Hamiltonian equations. The momentum matches the dimensionality of θ as a vector of length k .

We write the Hamiltonian function as $H(\theta, \mathbf{p})$, which consists of *potential* energy $U(\theta)$ and *kinetic* energy $K(\mathbf{p})$:

$$H(\theta, \mathbf{p}) = U(\theta) + K(\mathbf{p}), \mathbf{p}, \theta \in \mathbb{R}^k$$

In statistical applications of MCMC, we are primarily interested in generating θ from a given distribution $f(\theta)$. To do so, we let $U(\theta) = -\log f(\theta)$. For momentum, we typically assume $\mathbf{p} \sim N_k(0, \mathbf{M})$, where \mathbf{M} is a user-specified covariance matrix.

Hamiltonian Equations

Under the above formulation, we have that:

$$H(\boldsymbol{\theta}, \mathbf{p}) = -\log f(\boldsymbol{\theta}) + \frac{1}{2} \mathbf{p}' \mathbf{M} \mathbf{p}$$

Over time, HMC travels on trajectories that are governed by the following first-order differential equations, known as the Hamiltonian equations:

$$\begin{aligned} \frac{\partial \mathbf{p}}{\partial t} &= - \frac{\partial H(\boldsymbol{\theta}, \mathbf{p})}{\partial \boldsymbol{\theta}} = - \frac{\partial U(\boldsymbol{\theta})}{\partial \boldsymbol{\theta}} = \nabla_{\boldsymbol{\theta}} \log f(\boldsymbol{\theta}) \\ \frac{\partial \boldsymbol{\theta}}{\partial t} &= - \frac{\partial H(\boldsymbol{\theta}, \mathbf{p})}{\partial \mathbf{p}} = - \frac{\partial K(\mathbf{p})}{\partial \mathbf{p}} = \mathbf{M}^{-1} \mathbf{p} \end{aligned}$$

where $\nabla_{\boldsymbol{\theta}} \log f(\boldsymbol{\theta})$ is the gradient of the log posterior. A solution to the Hamiltonian equations is a function that defines the path of $(\boldsymbol{\theta}, \mathbf{p})$ from which specific values of $\boldsymbol{\theta}$ could be sampled.

Solving the Hamiltonian Differential Equations

- Solving the Hamiltonian equations is a critical step in HMC simulation.
- A standard approach for solving differential equations is Euler's method, which produces a discrete function that approximates the solution at each time t .
- Errors tend to accumulate in Euler's method, especially after a larger number of steps.
- In HMC, one often has to take a larger number of steps to ensure the new proposal is sufficiently far from the location of the previous sample.

Leapfrog method

The leapfrog method is a good alternative to the standard Euler's method for approximating the solutions to Hamiltonian equations (Ruth 1983).

The leapfrog algorithm modifies Euler's method by using a discrete step size ϵ individually for p and θ , with a full step ϵ in θ sandwiched between two half-steps $\epsilon/2$ for p :


$$p\left(t + \frac{\epsilon}{2}\right) = p(t) + \left(\frac{\epsilon}{2}\right) \nabla_{\theta} \log f(\boldsymbol{\theta}(t))$$

$$\theta(t + \epsilon) = \theta(t) + \epsilon M^{-1} p\left(t + \frac{\epsilon}{2}\right)$$

$$p(t + \epsilon) = p\left(t + \frac{\epsilon}{2}\right) + \left(\frac{\epsilon}{2}\right) \nabla_{\theta} \log f(\boldsymbol{\theta}(t + \epsilon))$$

Hamiltonian Monte Carlo

- Why does HMC work much better?
- Doesn't get stuck — follows gradient
- Extra variables (momentum, energy) provide diagnostics
- But also requires more
 - Gradients — curvature of log-posterior
 - “Mass” of particle
 - Number of leaps in a single trajectory
 - Size of individual leaps
- These need to be tuned right
- Gradients are unique to each model



The U-Turn Problem





Figure 9.6

- mc-stan.org
- Install RStan
 1. Get C++ compiler
 2. ???
 3. Profit



Interfaces


ways to run Stan

Stan Interfaces

The Stan modeling language and statistical algorithms are exposed through interfaces into many popular computing environments.

- RStan (R)
- PyStan (Python)
- CmdStan (shell, command-line terminal)
- MatlabStan (MATLAB)
- Stan.jl (Julia)
- StataStan (Stata)
- MathematicaStan (Mathematica)

Programs written in the Stan modeling language are portable across interfaces.



Stanislaw Ulam (1909–1984)



Check out the
movie:
“Adventures of a
Mathematician”

Stanislaw Ulam and his daughter Claire with MANIAC


Stan is NUTS




No U-Turn Sampler

Automatic Step Size and
Number Adaptation

- No U-Turn Sampler (NUTS2): Adaptive Hamiltonian Monte Carlo
- Implemented in Stan (rstan: mc-stan.org)





| | |
|-----------------------|---------------------------------------|
| Simulation options | |
| Algorithm | EfficientNUTS |
| Target distribution | donut |
| Autoplay | <input type="checkbox"/> |
| Autoplay delay | <input type="range" value="285"/> 285 |
| Tweening delay | <input type="range" value="18"/> 18 |
| Step | |
| Reset | |
| Visualization Options | |
| Algorithm Options | |
| Leapfrog Δt | <input type="range" value="0.1"/> 0.1 |
| Close Controls | |