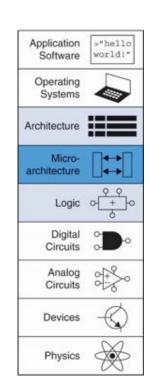
COSC175 (Systems I): Computer Organization & Design

Professor Lillian Pentecost Fall 2024

Warm-Up December 2

- Where we were
 - Multi-cycle vs. single cycle RISC-V processor
- Where we are going
 - Pipelined execution
- Logistics, Reminders
 - TA help 7-9PM on Sundays, Tuesdays, Thursdays in C107
 - Use to start review of whole semester!
 - LP Office hours M 9-10:30AM, Th 2:30-4PM
 - No more weekly exercises! Focus on your final lab, and starting to review previous assignments + exercises
 - BRING YOUR PROCESSOR LAB QUESTIONS INTO LAB WEDNESDAY!!
 - O ATTEND FACULTY CANDIDATE TALK TODAY AT 4:30PM in A131, WITH SNACKS @ 4PM in C209
 - If you attend at least 3 candidate talks, then send me an email including 1 thing you learned from each talk, I'll give you 5% extra credit on any previous lab report



Single- vs. Multicycle Processor

Single-cycle:

- + simple
- cycle time limited by longest instruction (lw)
- separate memories for instruction and data
- 3 adders/ALUs

• Multicycle: slower than single-cycle, unless clock cycle time and/or CPI improves

- + higher clock speed
- + simpler instructions run faster
- + reuse expensive hardware on multiple cycles
- sequencing overhead paid many times

Goals for Today

- Introduce alternative design: pipelined processor
- Keep in mind: what are the pros and cons of this design vs. single-cycle?
- How is the performance?

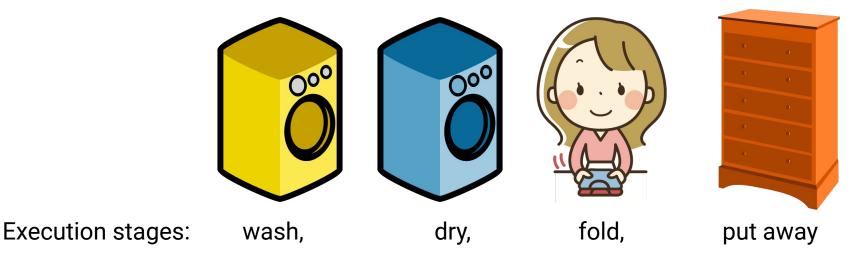
Same design steps as single-cycle:

- first datapath
- then control

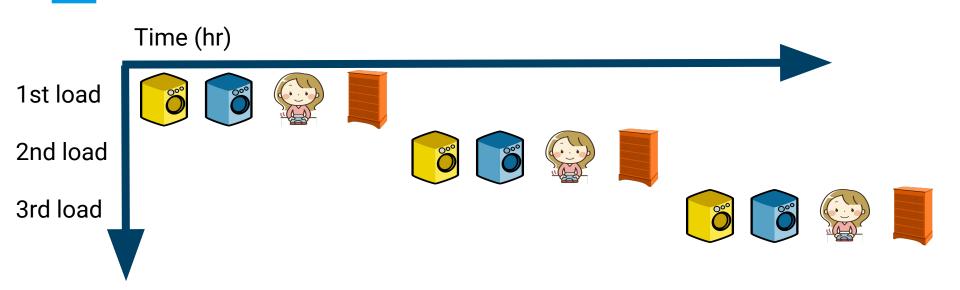
... remember the memory hierarchy, tho?

Introducing Pipelining

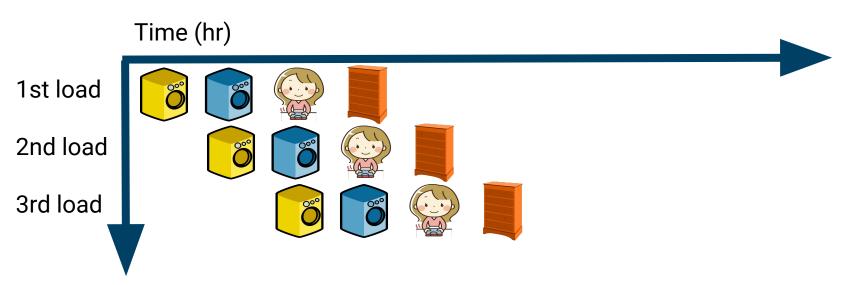
- Temporal Parallelism: multiple instructions in-progress at the same time
- Start with an analogy: loads of laundry
- HW units available: washer, dryer, folder (you & a table), dresser



Single-Cycle vs. Pipelined Laundry



Single-Cycle vs. **Pipelined** Laundry



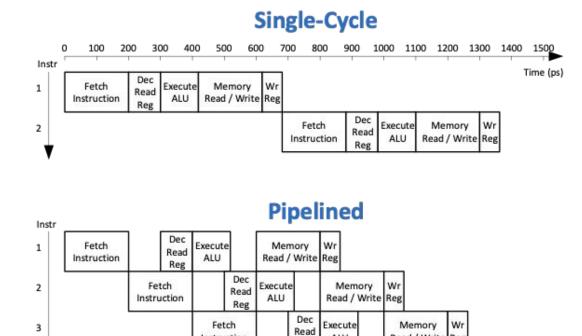
- Temporal Parallelism: multiple instructions in-progress at the same time
- HW units stay busy

Introducing a Pipelined Processor Design

- Temporal Parallelism: multiple instructions in-progress at the same time
- Divide single-cycle processor into 5 stages:
 - Fetch
 - Decode
 - Execute
 - Memory
 - Writeback
- Add pipeline registers between stages to propagate output of one stage as input to the next at the clock tick

Introducing a Pipelined Processor Design

- Looks just like our laundry! HW units are kept busy by starting execution of subsequent instructions
- What do we like about this idea?
- Do we have any questions or concerns about this design?

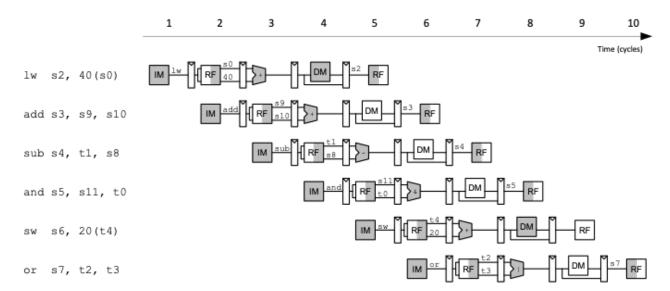


Read / Write Reg

Instruction

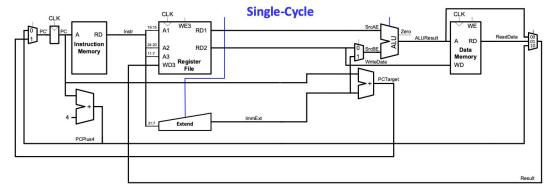
Introducing a Pipelined Processor Design

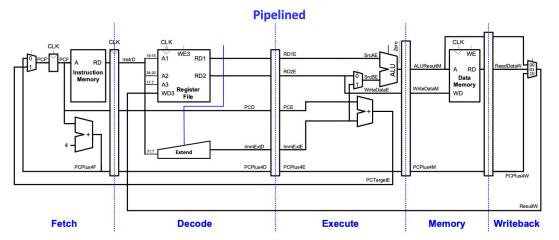
- Fetch reads IM, Decode reads RF, Execute uses ALU, Memory access DM, Writeback updates RF
- What do we like about this idea? What is your estimate of the average CPI?
- Do we have any questions or concerns about this design?



Updated Design Diagram

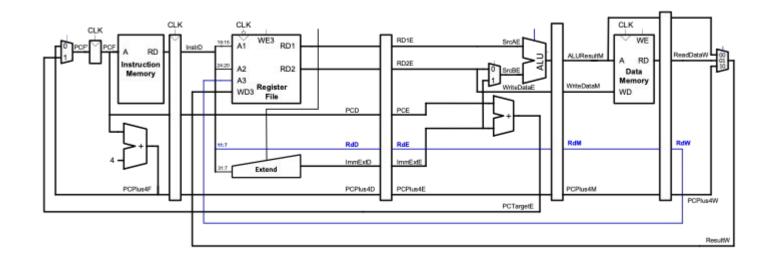
- Signals are now tied with specific stages, and we append the first letter of the relevant stage
- Are there any issues here?
- Let's trace through each stage





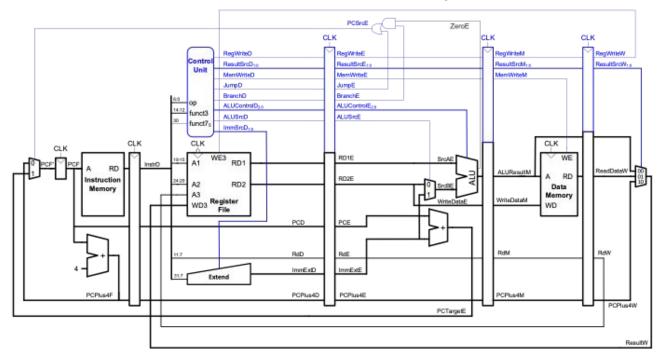
Updated Design Diagram: correct for WB

- Rd (destination register) should arrive at same time as Result for WB stage
- RF will be written on falling edge of CLK to account for delay



Control Unit is Same as Single-Cycle

Control Signals travel with instruction stage-by-stage



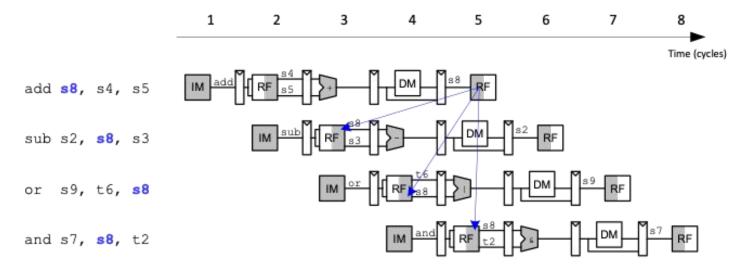
Let's use the rest of class to discuss

 The Good News: performance of the pipelined processor is better than the single-cycle (we'll compute by how much)

- The Bad News: We introduce hazards subsequent instructions may depend on an instruction that hasn't completed yet
 - Data Hazard if a required register value is not yet written back to the register file
 - Control Hazard if the next instruction to execute is not determined yet (i.e., branch if X)

Data Hazard Example

sub relies on the result of add, or relies on the result of sub, and so on...



What might we do to fix this?

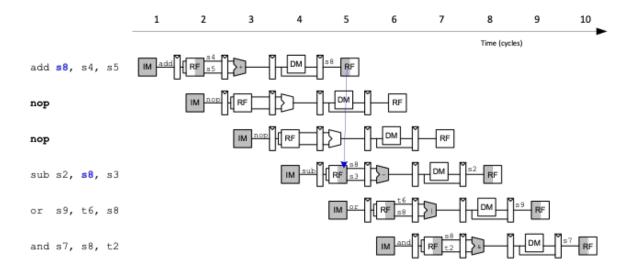
Resolving Data Hazards

- At the time we compile our code, we might check for data hazards and choose to:
 - Rearrange instruction order
 - Insert nops between dependent instructions

- At execution time (or run time), we might choose to:
 - Stall (i.e., sit idle) until the required value is ready
 - "Forward" data to a dependent instruction as soon as it is ready

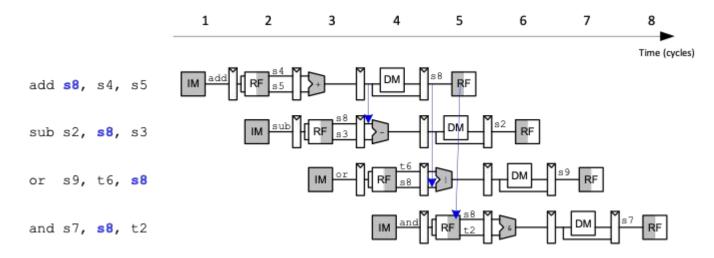
Resolving Data Hazards: Compile Time

- Fill the time with useful, independent instructions until hazard resolves
- If there are no such independent instructions, insert nop



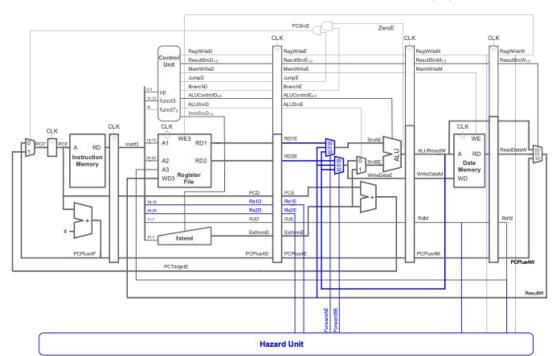
Resolving Data Hazards: Run Time

- Value needed will exist (i.e., ALU produced it) before it is written back to the RF
- Data Forwarding is a "sneak peak" of a previous instruction's result to avoid stalling
- Compare if an instruction's source register in Execute stage matches the destination register of an instruction in the Memory or WB stage; if yes, then forward ALUResult



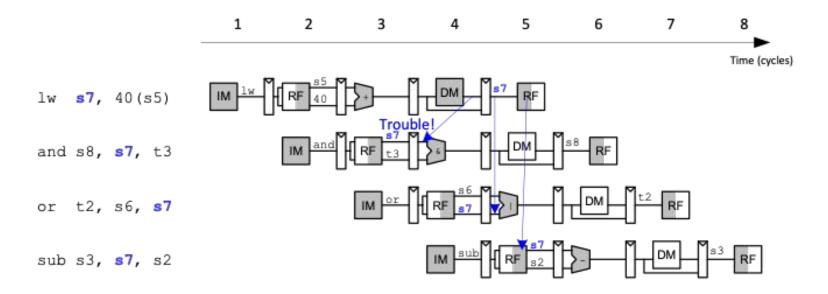
Resolving Data Hazards: Run Time

"Hazard Unit" will check hazard condition(s) each clock cycle

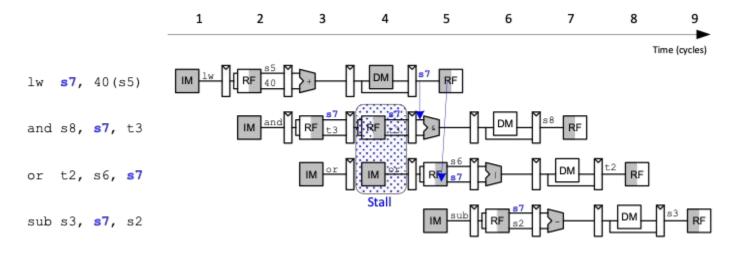


Resolving Data Hazards

What other data hazards may exist?



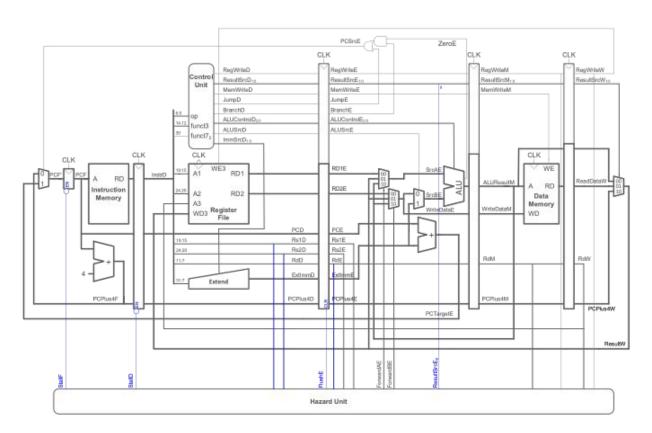
Resolving Data Hazards: sometimes you just have to be patient (stall)



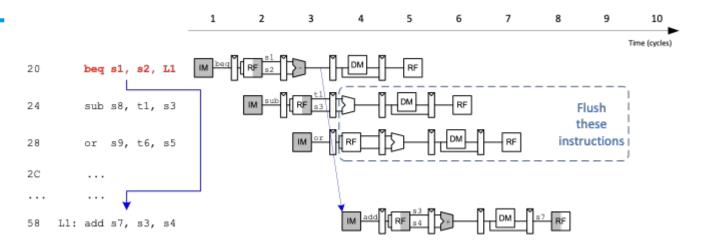
- Stall condition:
 - Is either **source** register in the Decode stage the same as the **destination** register in the Execute stage?

 AND
 - o Is the instruction in the Execute stage a lw?

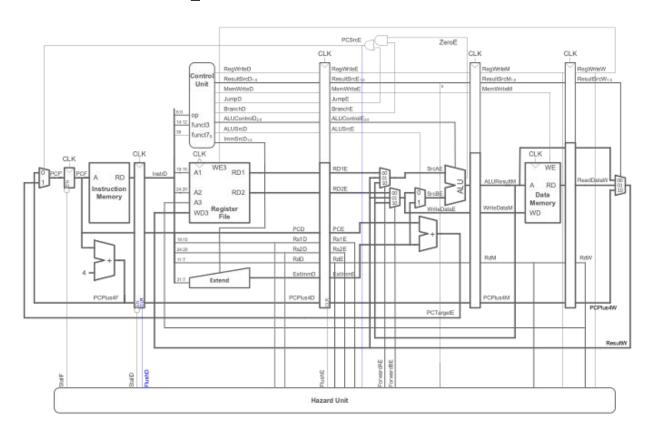
Pipelined Processor with Data Hazard Protection



What about **control** hazards?



- We don't determine whether the branch is taken until execute stage
- If the branch is taken, we've started execution of 2 instructions that we shouldn't have! Uh oh!!
- In the case of a taken branch, we must flush the in-progress instructions and start fresh – in this case, a branch misprediction penalty of 2 instructions



- Remember SPECINT2000 benchmarks:
 - 25% loads
 - 10% stores
 - 13% branches
 - 52% R-type
- Suppose:
 - 40% of loads used by next instruction (causes 1 cycle stall)
 - 50% of branches mispredicted (causes 2 cycle penalty)
- What do we need to compute the performance?
 - CPI and Clock Cycle time, as usual!

- Remember SPECINT2000 benchmarks:
 - o 25% loads
 - o 10% stores
 - 13% branches
 - 52% R-type
- Suppose:
 - 40% of loads used by next instruction (causes 1 cycle stall)
 - 50% of branches mispredicted (causes 2 cycle penalty)
- Average CPI: unfortunately, it isn't 1
 - CPI for a load instruction is 2 cycles 40% of the time
 - CPI for a branch is 3 cycles 50% of the time

Average CPI = (% load) (cycles per load) + (% store) (cycles per store) + (% branch) (cycles per branch) + (% R-type) (cycles per R-type)

Average CPI = (0.25)(1*0.6+2*0.4) + (0.1)(1) + (0.13)(2) + (0.52)(1) =**1.23**

Clock Cycle time of pipelined design? Determine critical path!

```
Tc_pipelined = max of

(tpcq + tmem + tsetup) Fetch

2(tRFread + tsetup) Decode

(tpcq + 4tmux + tALU + tAND-OR + tsetup)Execute

(tpcq + tmem + tsetup) Memory

2(tpcq + tmux + tRFwrite) Writeback
```

Note: the 2X is because Decode and Writeback stages both use the register file in each cycle — Each stage gets half of the cycle time (Tc/2) to do their work, or, stated a different way, 2x of their work must fit in 1 cycle

Clock Cycle time of pipelined design? Determine critical path!

```
Tc_pipelined = max of

(tpcq + tmem + tsetup) Fetch

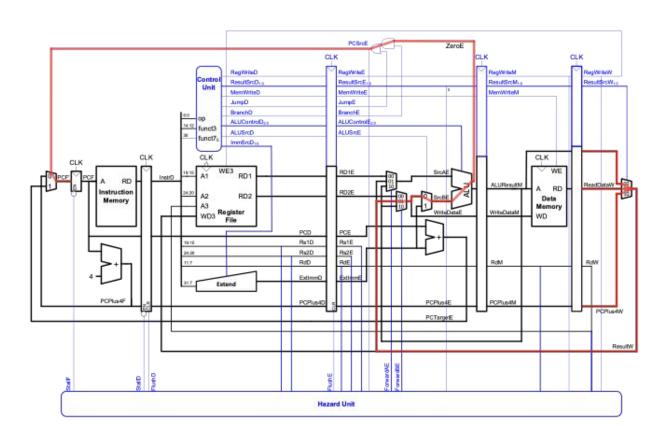
2(tRFread + tsetup) Decode

(tpcq + 4tmux + tALU + tAND-OR + tsetup) Execute is critical path

(tpcq + tmem + tsetup) Memory

2(tpcq + tmux + tRFwrite) Writeback
```

Note: the 2X is because Decode and Writeback stages both use the register file in each cycle — Each stage gets half of the cycle time (Tc/2) to do their work, or, stated a different way, 2x of their work must fit in 1 cycle



Element	Parameter	Delay (ps)
Register clock-to-Q	$t_{pcq_{ m PC}}$	40
Register setup	$t_{ m setup}$	50
Multiplexer	$t_{ m mux}$	30
AND-OR gate	$t_{ m AND-OR}$	20
ALU	$t_{ m ALU}$	120
Decoder (Control Unit)	$t_{ m dec}$	25
Extend unit	$t_{ m dec}$	35
Memory read	$t_{ m mem}$	200
Register file read	$t_{RF\mathrm{read}}$	100
Register file setup	t_{RF} setup	60

$$T_{c_pipelined} = t_{pcq} + 4t_{mux} + t_{ALU} + t_{AND-OR} + t_{setup}$$

= $(40 + 4*30 + 120 + 20 + 50) \text{ ps} = 350 \text{ ps}$

We have an estimated CPI, a clock cycle time, let's consider our example program!

Check-in Activity, in pairs: What is the program execution time of a program with 100 billion instructions with our pipelined design (assume SPEC CPI)? How much faster than single-cycle?

Repeat for extreme cases:

- (1) If we **always** mispredict branches and we **always** have a dependency following 1w in this case, does pipelined design do better or worse than single cycle?
- (2) If we **never** mispredict branches and we **never** have a dependency following 1w in this case, does pipelined design do better or worse than single cycle?

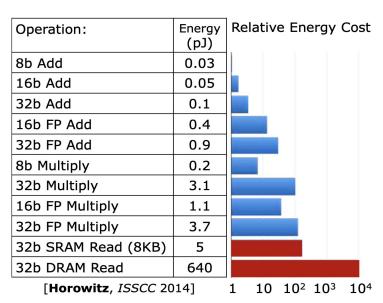
Computer design question to worry about over break

We have been using an ideal (super unrealistic) memory array

How can we maintain program performance if memory access actually

takes at least 10-100x longer than our ALU?

	Capacity	Read Latency (approx)
RegFile	< 1KB	10-100ps
SRAM	1-64MB	1-10ns
DRAM	4-64GB	~100ns (maybe more)
Flash, SSD, hard drive	256-1024GB	~100us



Wrap-Up December 2



- Coming up next!
 - Realistic Memory Systems; looking ahead to COSC275
- Logistics, Reminders
 - TA help 7-9PM on Sundays, Tuesdays, Thursdays in C107
 - LP Office hours M 9-10:30AM, Th 2:30-4PM
 - After receiving additional requests, weekly exercise solutions will be posted next week (Dec
 8) to help with final exam review
 - BRING YOUR PROCESSOR LAB QUESTIONS INTO LAB WEDNESDAY!!
 - Final Lab Report (all stages) due Monday, Dec. 9 10PM
 - O ATTEND FACULTY CANDIDATE TALK TODAY AT 4:30PM in A131, WITH SNACKS @ 4PM in C209
 - If you attend at least 3 candidate talks, then send me an email including 1 thing you learned from each talk, I'll give you 5% extra credit on any previous lab report
 - POLL FOR ADVANCED TOPIC!! Vote tonight if you haven't already!!!
- FEEDBACK
 - https://forms.gle/5Aafcm3iJthX78jx6